

Istio – An introduction for developers

WW Developer Advocate Team

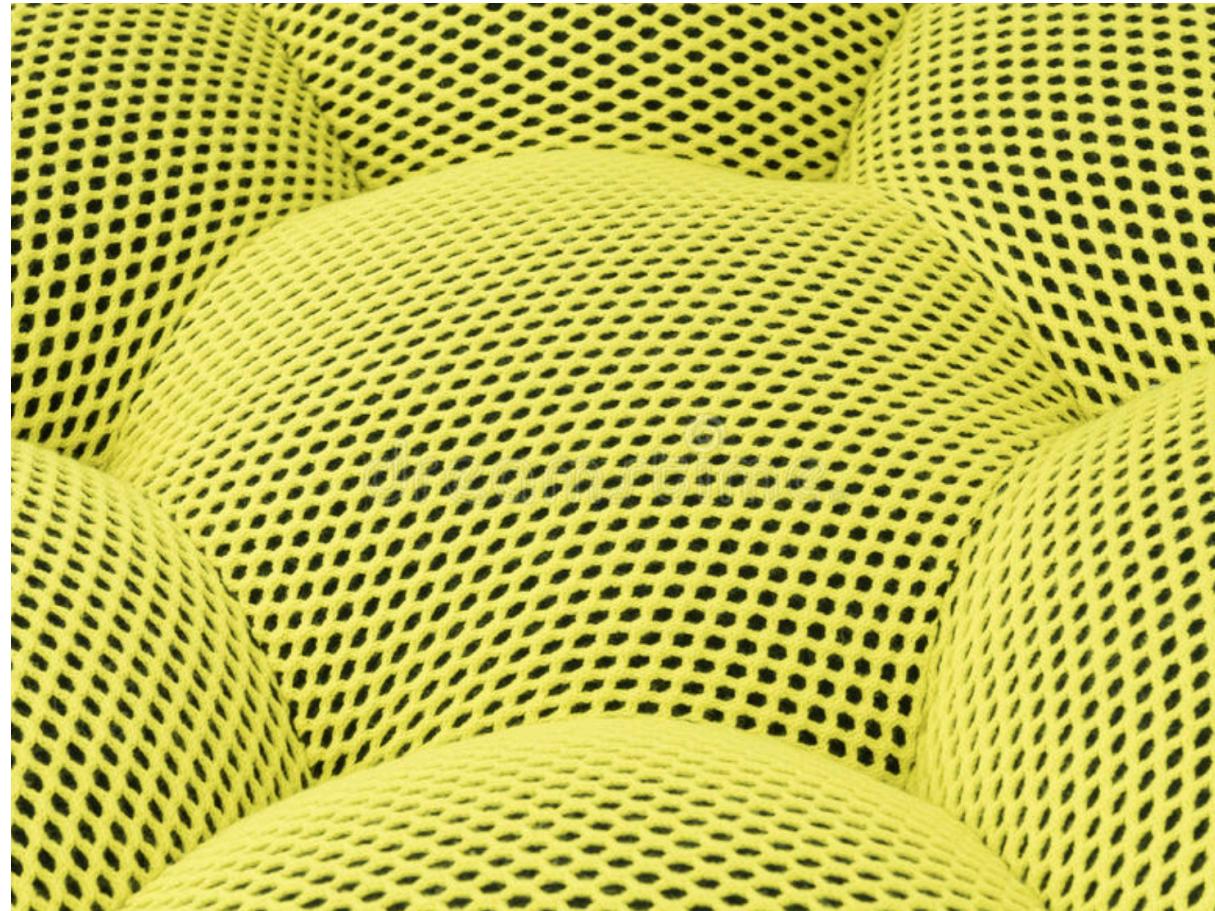
IBM Developer

Istio - An open platform that provides a uniform way to **connect, manage, and secure microservices**.

Istio supports managing traffic flows between microservices, enforcing access policies, and aggregating telemetry data, all **without requiring changes to the microservice code**.

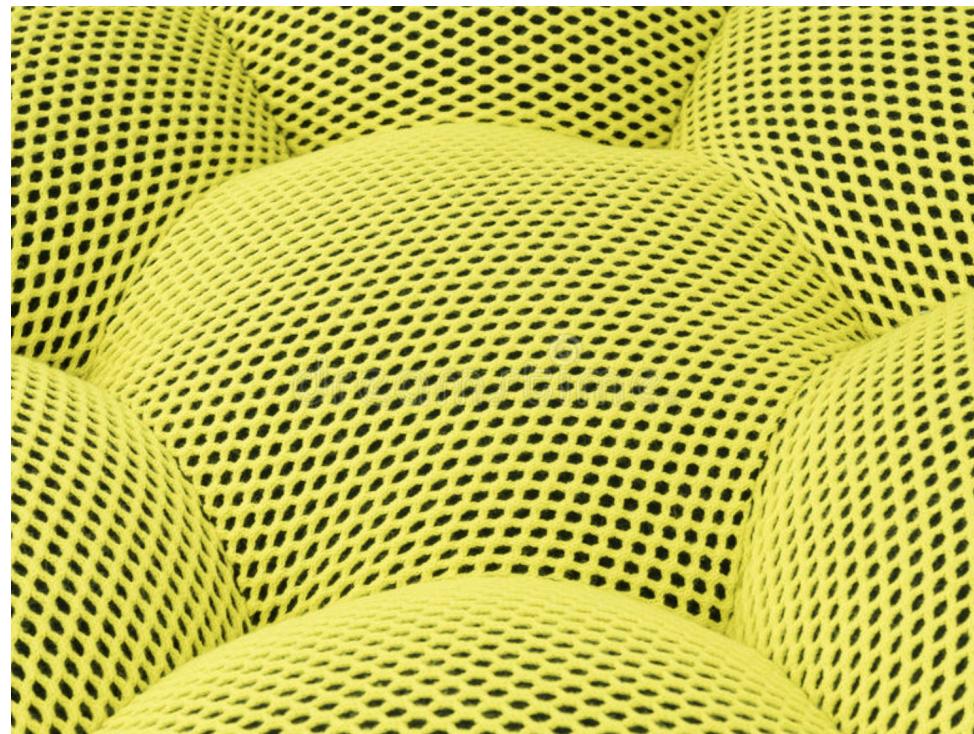
Service Mesh & Istio Architecture

What is a service mesh?



A service mesh provides a **transparent** and **language-independent** way to flexibly and easily manage the **communication** between microservices.

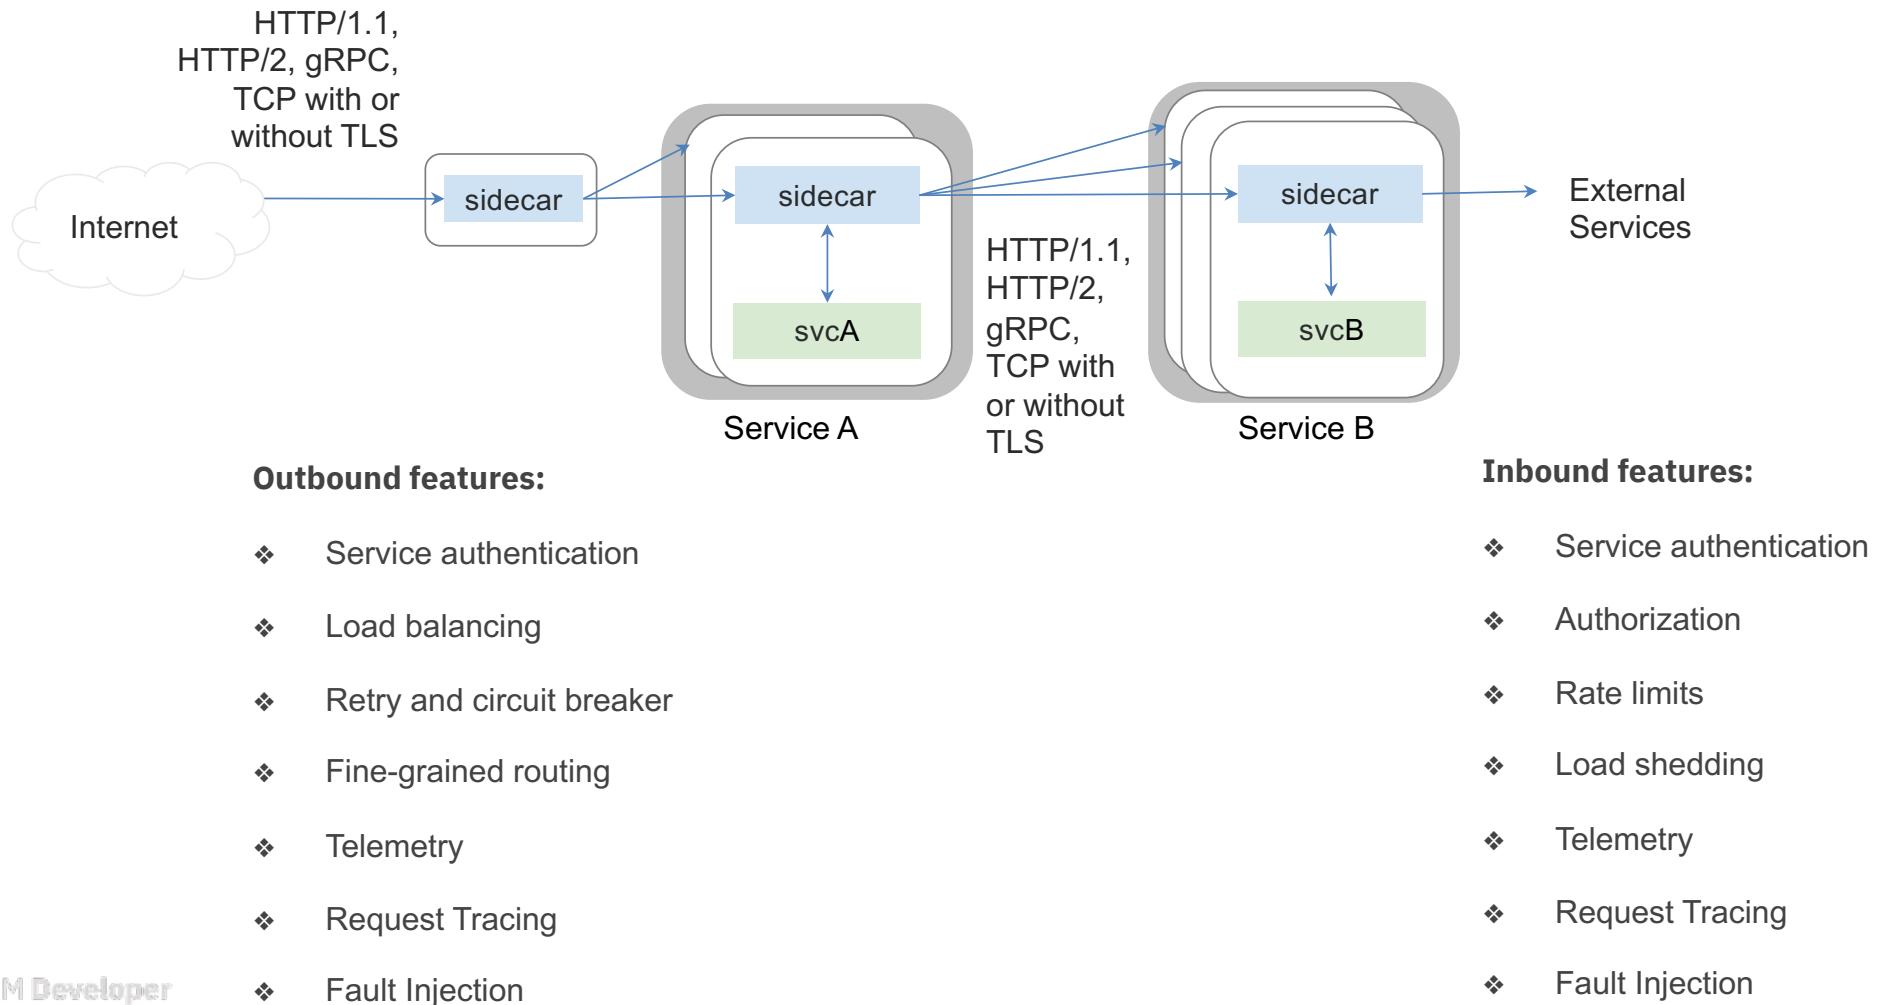
What is a service mesh?



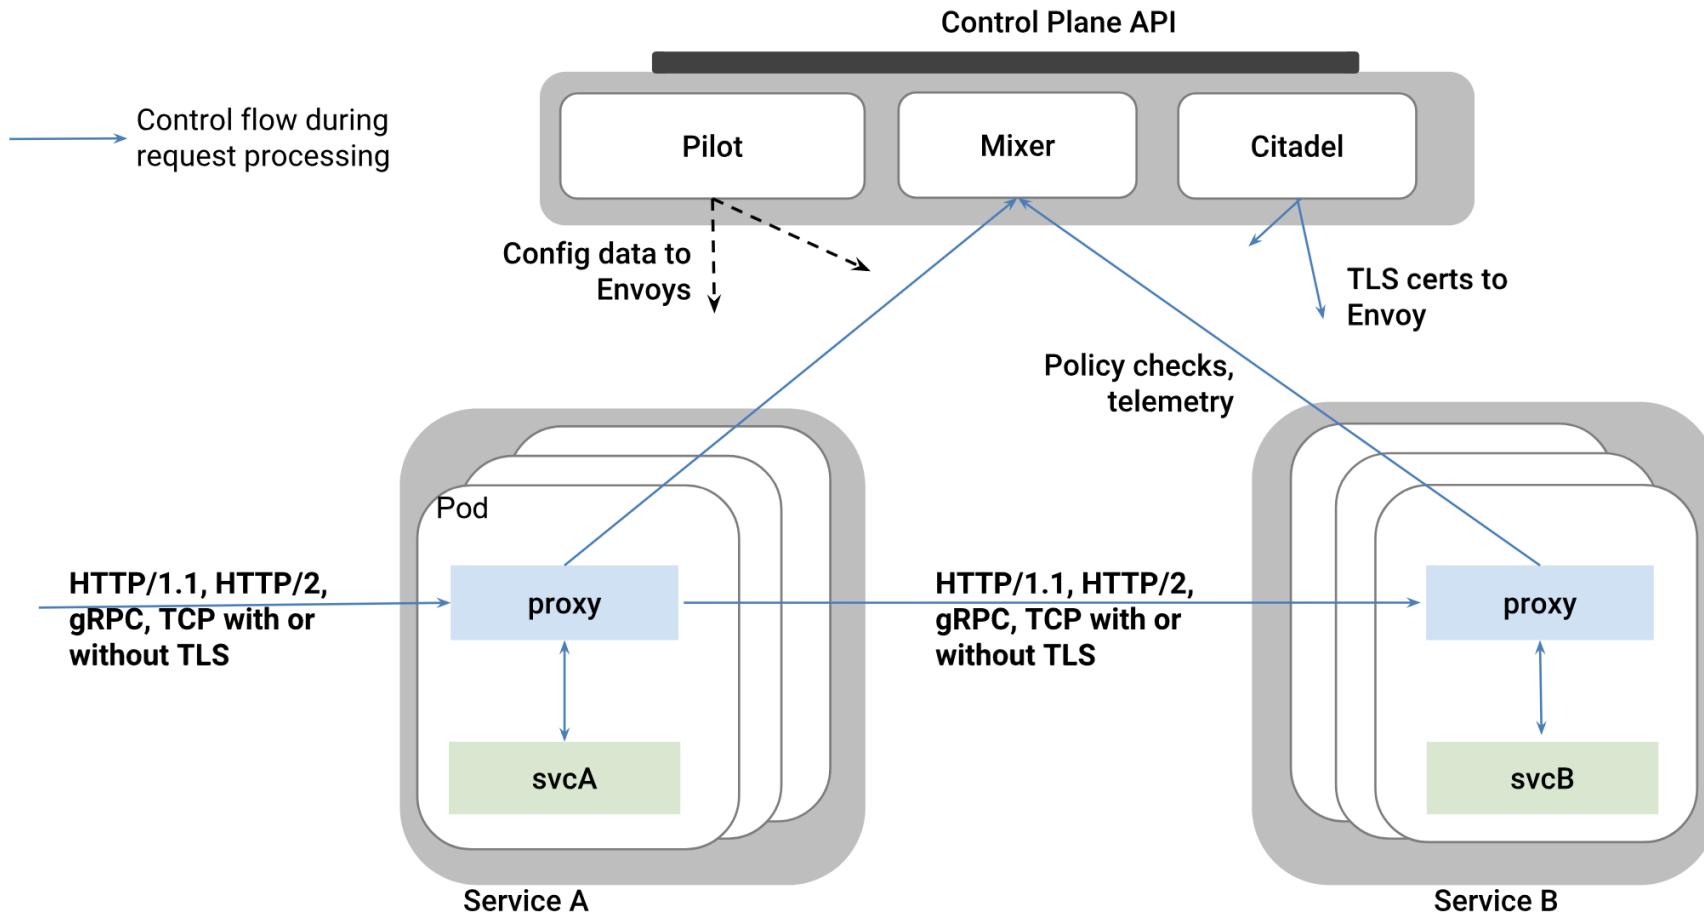
Service meshes can assist you with microservice implementations:

- How can I **find** the service I need?
- How can I **scale** my services?
- How can I test new **versions** of services without impacting new users?
- How can I test against **failures**?
- How can I **secure** service-to-service communication?
- How can I **route** traffic in a specific way?
- How can I test **circuit breaking** or fault injection?
- How can I **monitor** my microservices and collect **metrics**?
- How can I do **tracing**?
- Most importantly, how can I do all these things **without changing individual microservices application **Code**?**

Weaving the mesh



Istio Architecture



Istio Architecture

Mixer

Mixer is a platform-independent component responsible for enforcing access control and usage policies across the service mesh and collecting telemetry data from the [Envoy proxy](#) and other services. The proxy extracts request level attributes, which are sent to Mixer for evaluation.

Pilot

Pilot provides service discovery for the Envoy sidecars, traffic management capabilities for intelligent routing (e.g., A/B tests, canary deployments, etc.), and resiliency (timeouts, retries, circuit breakers, etc.). It converts high level routing rules that control traffic behavior into Envoy-specific configurations, and propagates them to the sidecars at runtime.

Citadel

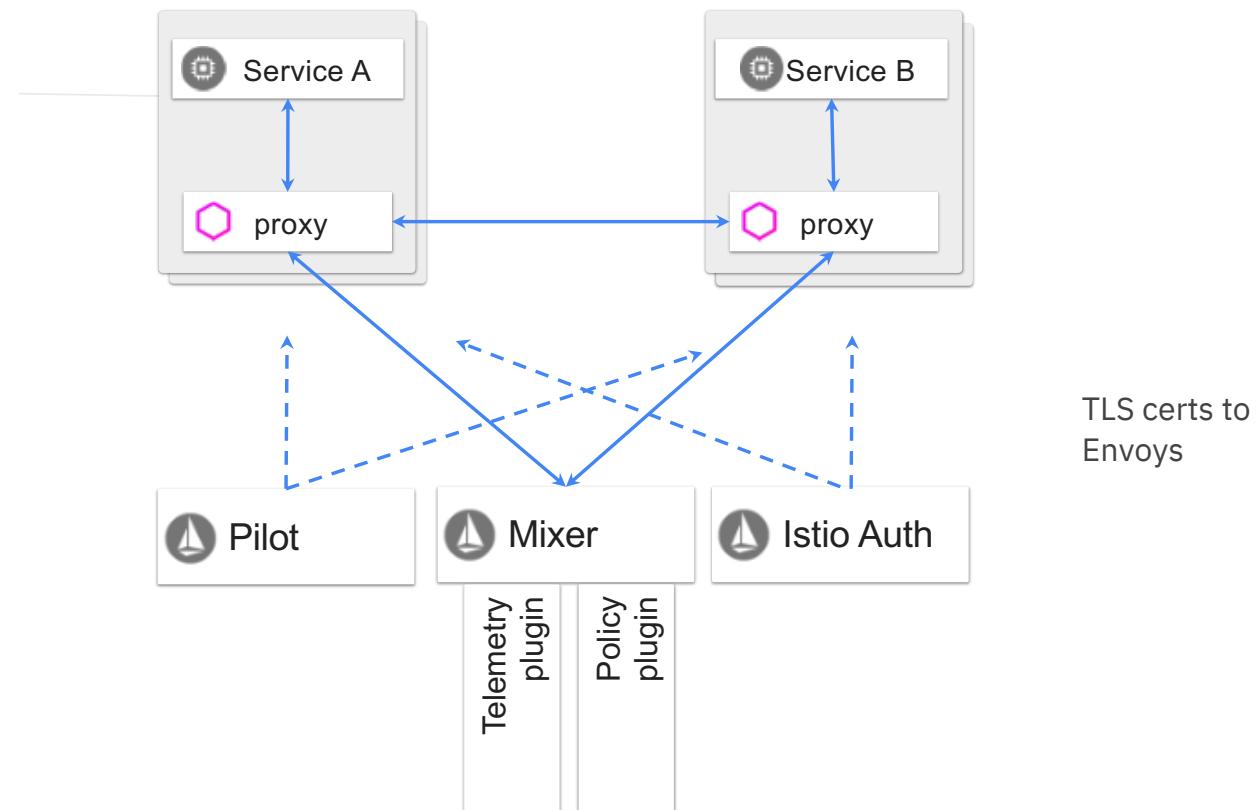
Citadel provides strong service-to-service and end-user authentication, with built-in identity and credential management.

Istio Architecture

Putting it Together

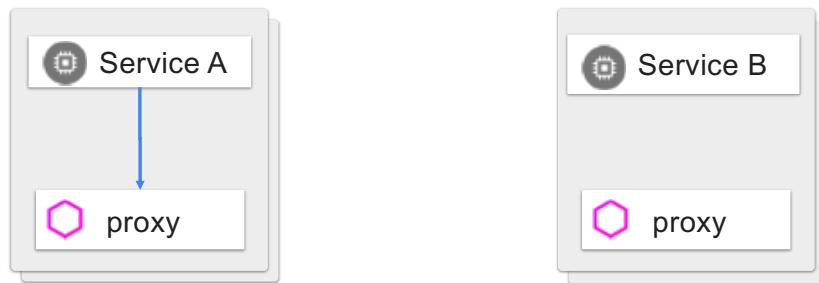
Traffic is transparently intercepted and proxied.
(App is unaware of Proxy presence)

Routing
and load
balancing
config to
Envoy

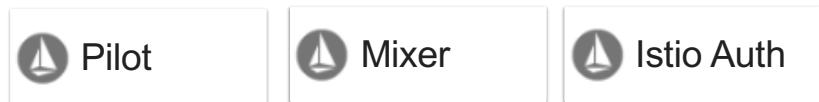


TLS certs to
Envoy

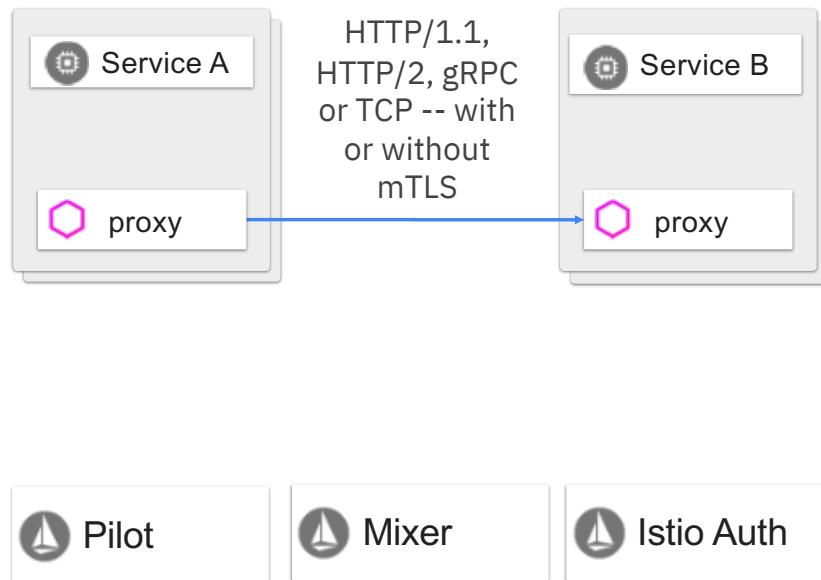
Life of a request



1. Service A places a call to <http://service-b>.
2. Client-side Envoy intercepts the call.
3. Envoy consults config (previously received from Pilot) to know how/where to route call to service B (taking into account service discovery, load balancing, and routing rules),
4. forwards the call to the right destination.

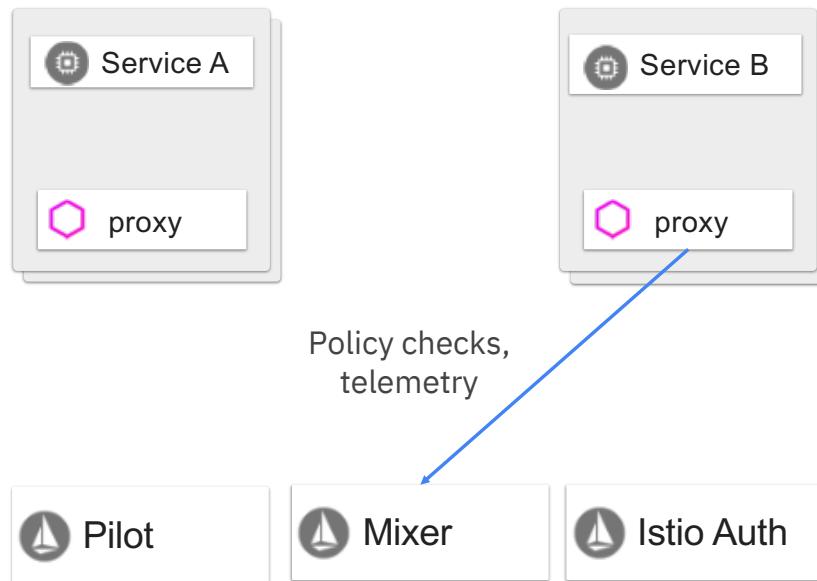


Life of a request



Envoy (A) forwards request to appropriate instance of service B. **Envoy** (B) proxy deployed with the service intercepts the call.

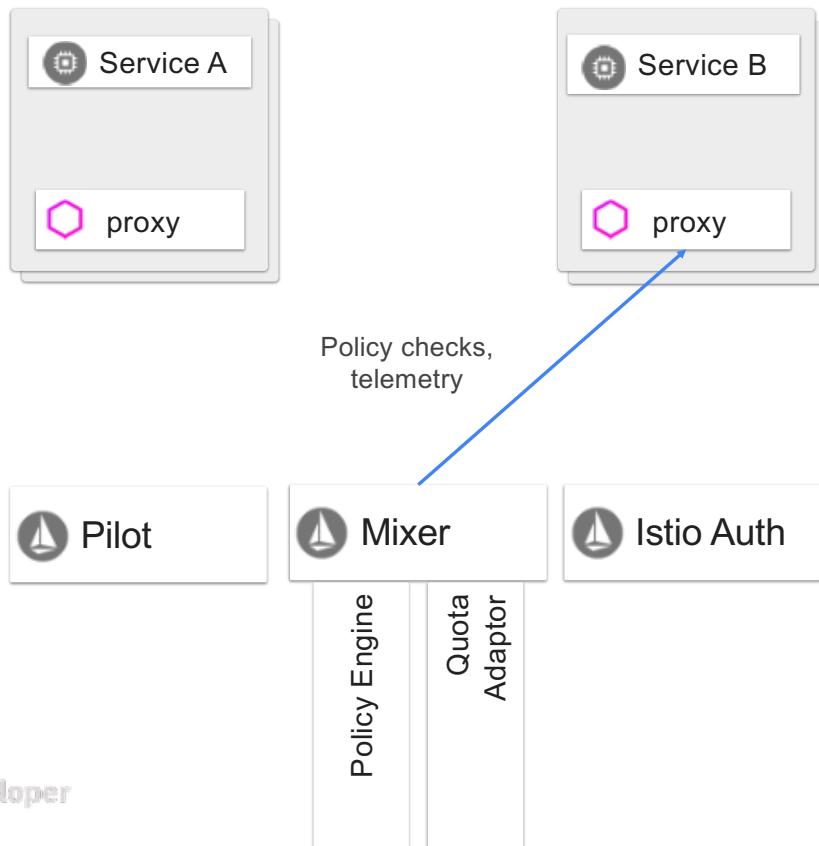
Life of a request



Server-B Envoy checks with Mixer to validate that call should be allowed

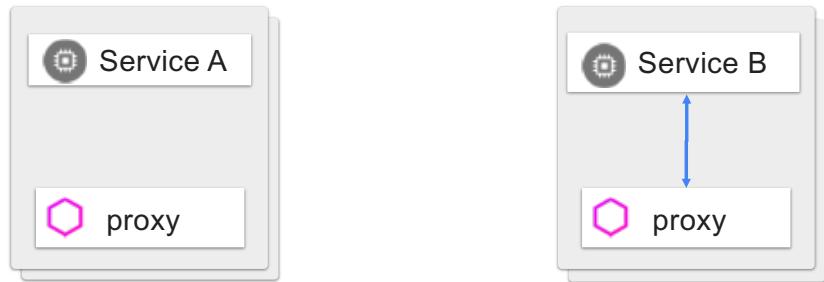
- ACL check
- quota check
- etc.

Life of a request

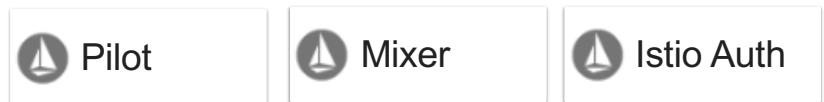


Mixer checks with appropriate adapters (policy engine, quota adapter) to verify that the call can proceed and returns true/false to Envoy

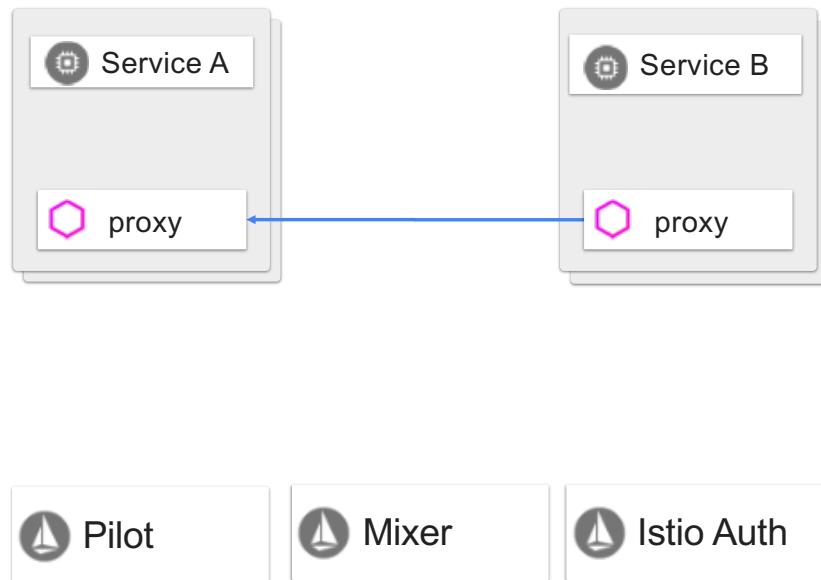
Life of a request



Server-side Envoy forwards request to service B, which processes request and returns response

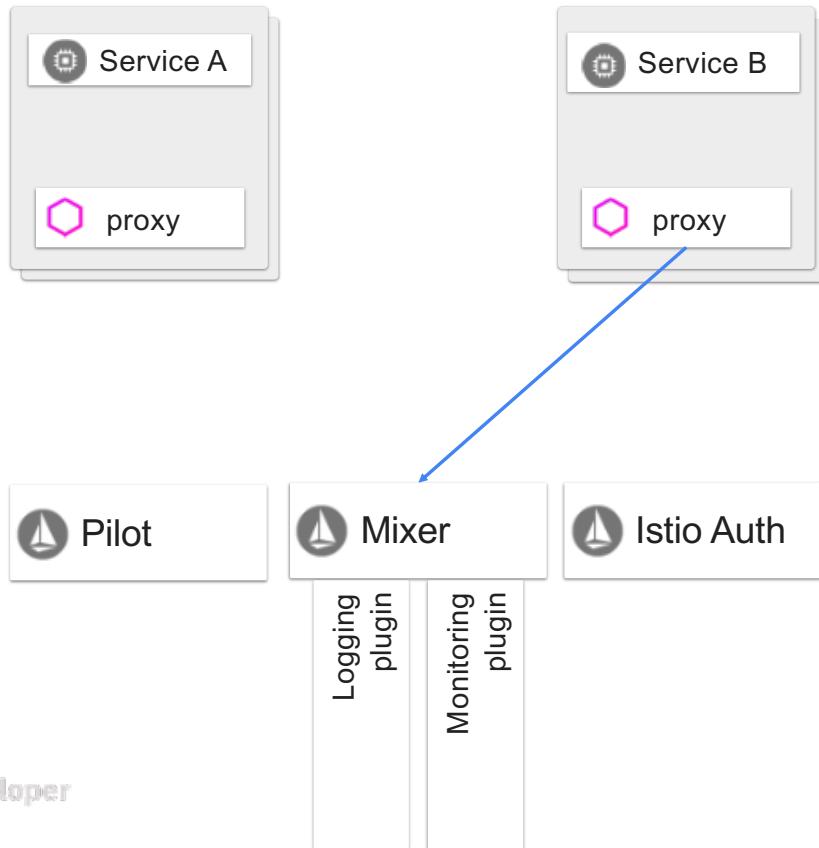


Life of a request



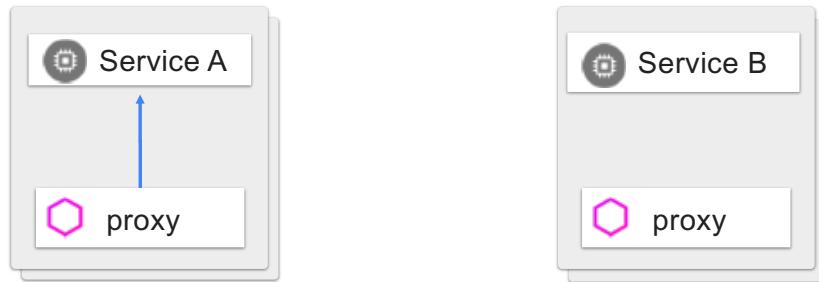
Envoy forwards response to the original caller, where response is intercepted by Envoy on the caller side.

Life of a request

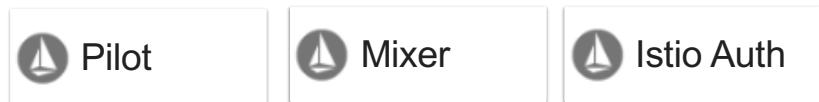


Envoy reports telemetry to Mixer, which in turn notifies appropriate plugins

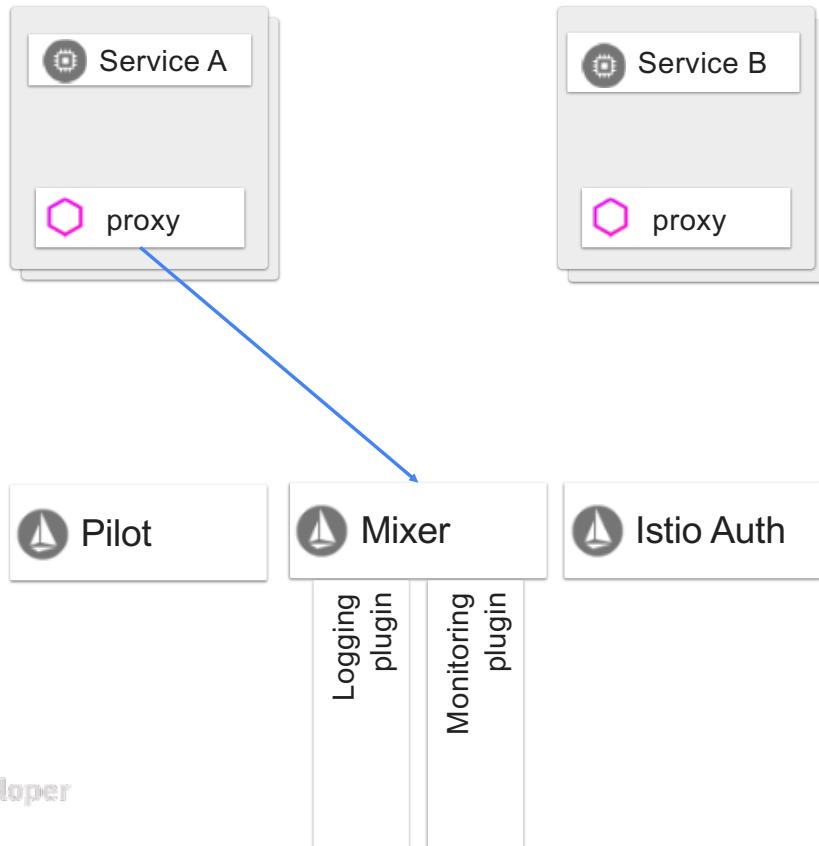
Life of a request



Client-side Envoy forwards response to original caller.



Life of a request



Client-side Envoy reports telemetry to Mixer (including client-perceived latency), which in turn notifies appropriate plugins

Istio Proxy Sidecar

To create a service mesh with Istio, you update the deployment of the pods to add the Istio Proxy (based on the [Lyft Envoy Proxy](#)) as a side car to each pod.

The Proxy then run as a separate container that manages all communication with that pod.

Istio modifies the individual pods and not the deployments

Manual Approach

To create a service mesh with Istio, you update the deployment of the pods to add the Istio Proxy (based on the Lyft Envoy Proxy) as a side car to each pod.

```
istioctl kube-inject -f 03hub-deployment.yaml
```

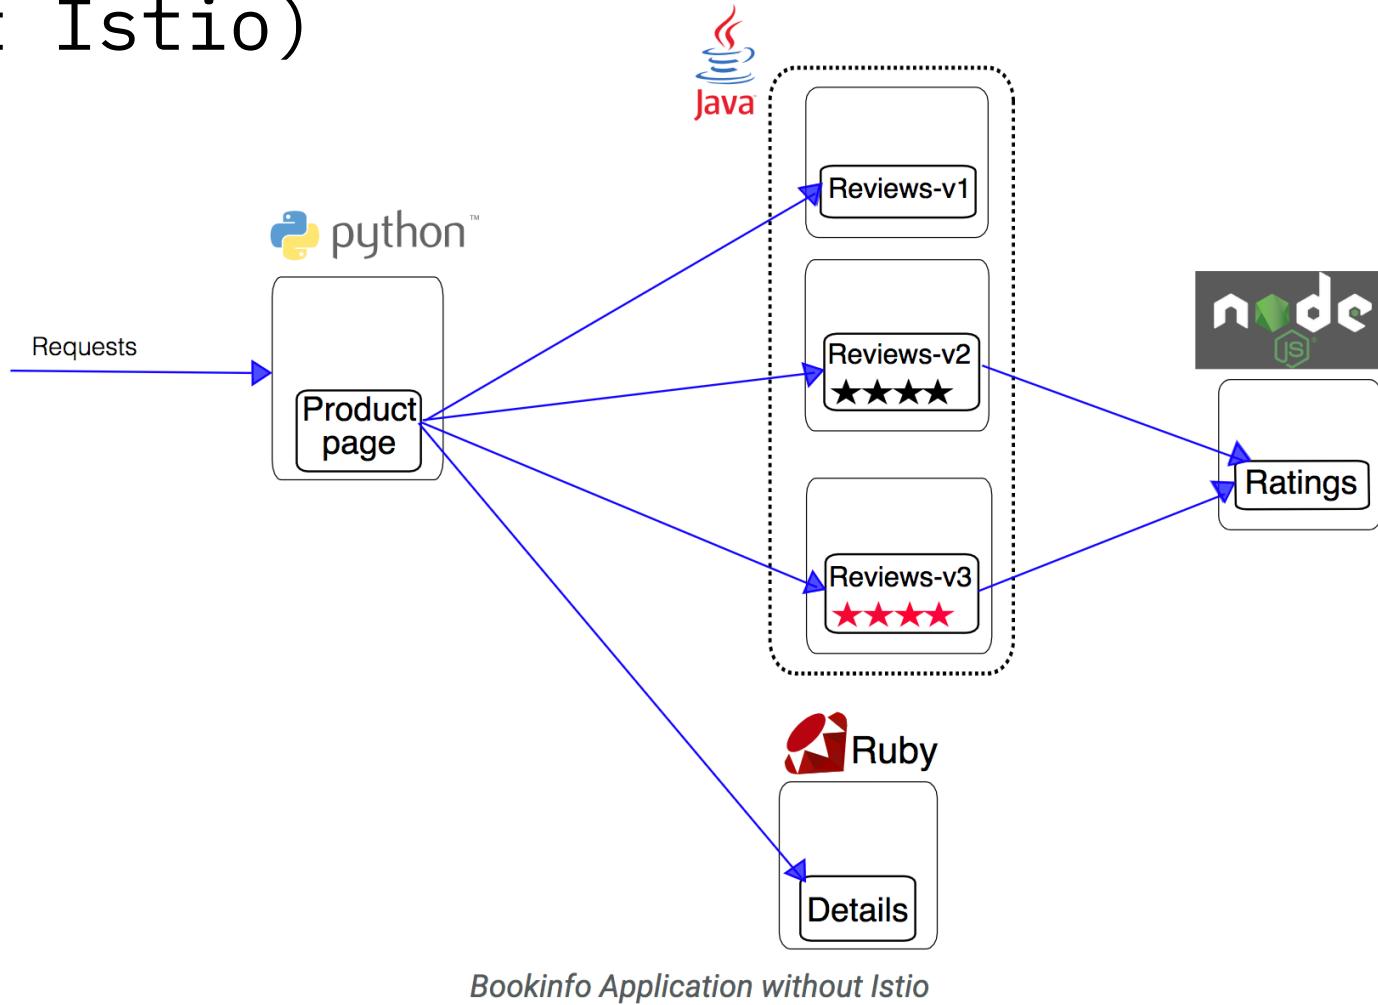
```
image: docker.io/istio/proxyv2:0.8.0
imagePullPolicy: IfNotPresent
name: istio-proxy
resources:
  requests:
    cpu: 100m
    memory: 128Mi
  securityContext:
```

Automatic Approach

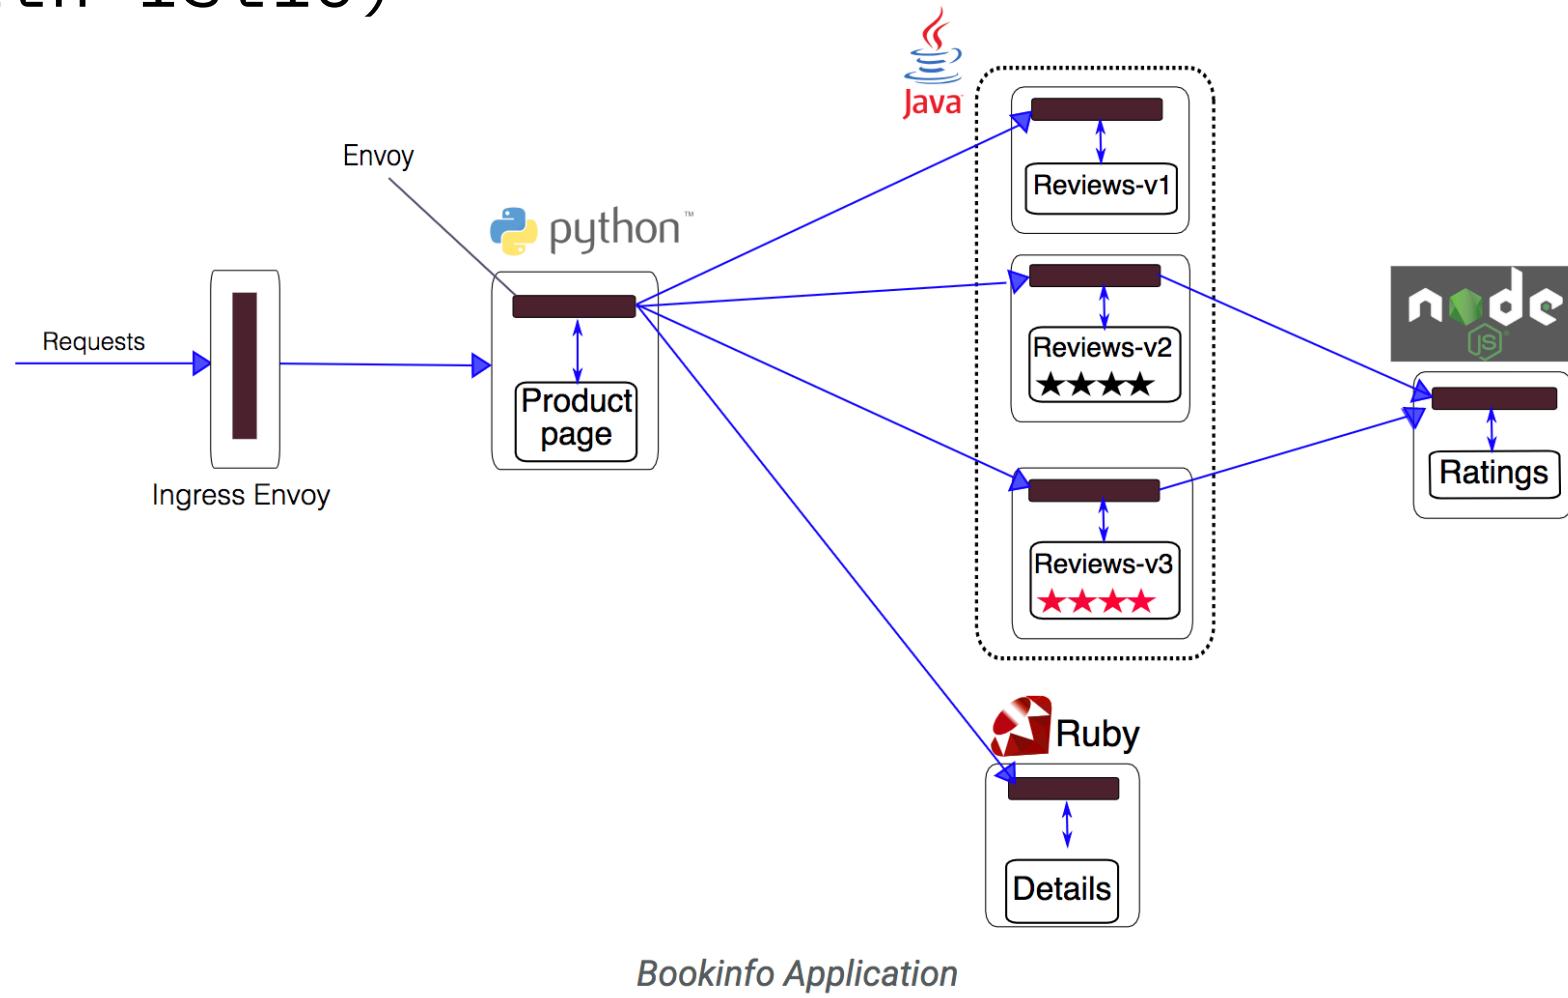
- Istio sidecars can also be automatically injected into a pod at creation time using a feature in Kubernetes called a [mutating webhook admission controller](#).
- When you install istio, a deployment is automatically created called istio-sidecar-injector-* that injects proxy side-cars automatically to pods installed on namespaces that have a label istio-injection=enabled

```
kubectl label namespace default istio-injection=enabled
```

Sample Microservices Application (without Istio)



Sample Microservices Application (with Istio)



Istio Ingress Controller

Control Ingress Traffic

Gateway describes a load balancer operating at the edge of the mesh receiving incoming or outgoing HTTP/TCP connections.

VirtualService defines a set of traffic routing rules to apply when a host is addressed. Each routing rule defines matching criteria for traffic of a specific protocol.

Istio ingress controller gets exposed as a normal **kubernetes service load balancer on port 80**.

```
kubectl get svc -n istio-system
```

<https://istio.io/docs/tasks/traffic-management/ingress/>

Traffic Management & Canary Deployments

Core Components

The core component used for traffic management in Istio is Pilot, which manages and configures all the Envoy proxy instances deployed in a particular Istio service mesh.

VirtualService defines a set of traffic routing rules to apply when a host is addressed. Each routing rule defines matching criteria for traffic of a specific protocol. If the traffic is matched, then it is sent to a named destination service (or subset or version of it) defined in **DestinationRule**.

DestinationRule defines policies that apply to traffic intended for a service after routing has occurred. Any destination host and subset referenced in a VirtualService rule must be defined in a corresponding DestinationRule.

ServiceEntry configuration enables services within the mesh to access a service not necessarily managed by Istio. The rule describes the endpoints, ports and protocols of a white-listed set of mesh-external domains and IP blocks that services in the mesh are allowed to access. For example for accessing a Watson service.

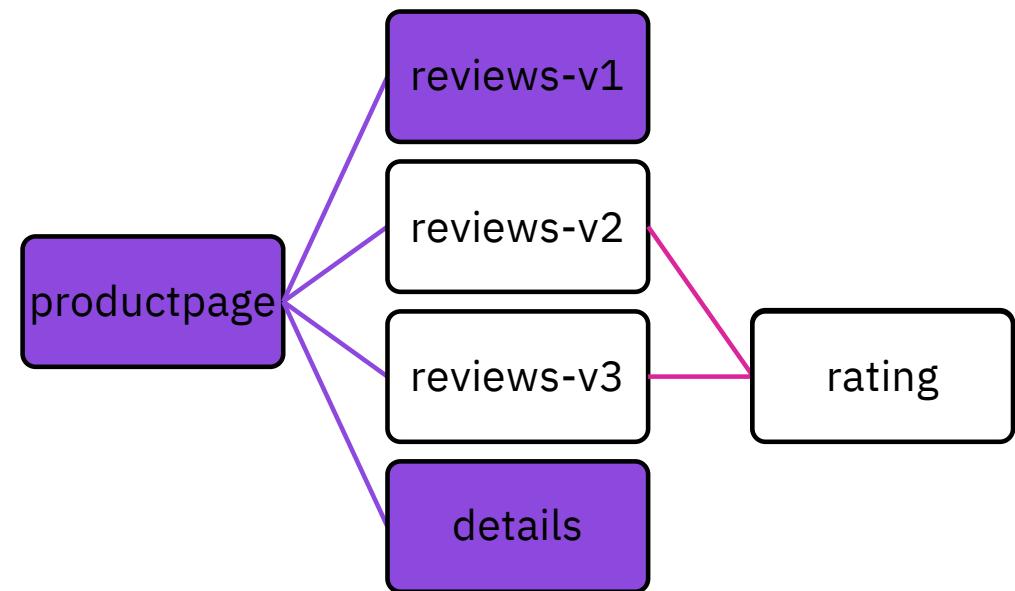
<https://istio.io/docs/reference/config/istio.networking.v1alpha3/>

Possible Use Cases

- A/B Testing
- Canary Deployments
- Route based on header
- Fault Injection
- HTTP Retry
- Rate Limits/Polices
- Circuit Breaker

A/B Testing

- A/B testing is a method of performing identical tests against two separate service versions in order to determine which performs better.
- Enforce only v1 of reviews microservice



A/B Testing

A/B testing is a method of performing identical tests against two separate service versions in order to determine which performs better.

Enforce only v1 of reviews microservice

```
istioctl create -f 02route-rule-reviews-v1.yaml
```

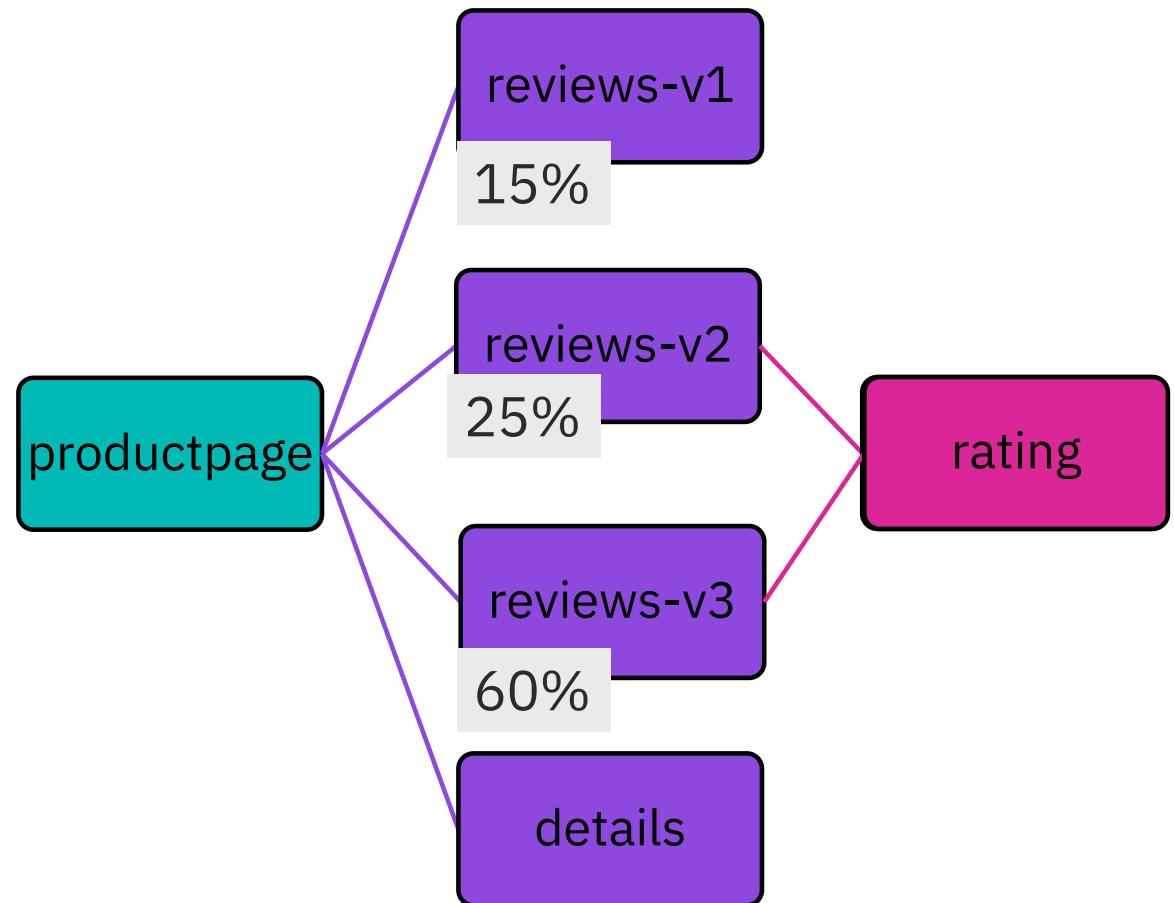
```
02route-rule-reviews-v1.yaml
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
  - reviews
  http:
  - route:
    - destination:
      host: reviews
      subset: v1
---
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  subsets:
  - name: v1
    labels:
      version: v1
```

Canary Deployments

- You can do Canary Deployments with Kubernetes replicaset alone, but it's quite complex.
- Maintaining canary traffic at small percentages requires many replicas (e.g., 1% would require a minimum of 100 replicas).
- Istio's routing rules also provide other important advantages; you can easily control fine grain traffic percentages (e.g., route 1% of traffic without requiring 100 pods) and you can control traffic using other criteria (e.g., route traffic for specific users to the canary version).
- All the traffic inside the pod is captured by iptables commands and directed to the sidecar proxy. Then the sidecar proxy performs routing, according to routing tables it receives from Istio Pilot (a part of the Istio Control Plane)

Canary Deployments

- Enforce the following:
- 15% for reviews-v1
- 25% for reviews-v2
- 60% for reviews-v3

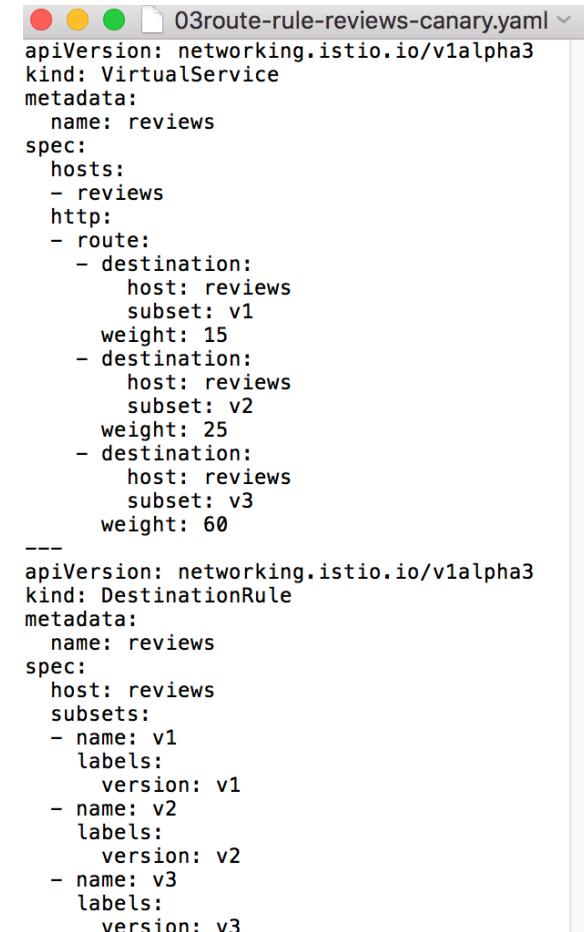


Canary Deployments

Enforce the following:

- 15% for reviews-v1
- 25% for reviews-v2
- 60% for reviews-v3

```
istioctl replace -f 03route-rule-reviews-
canary.yaml
```

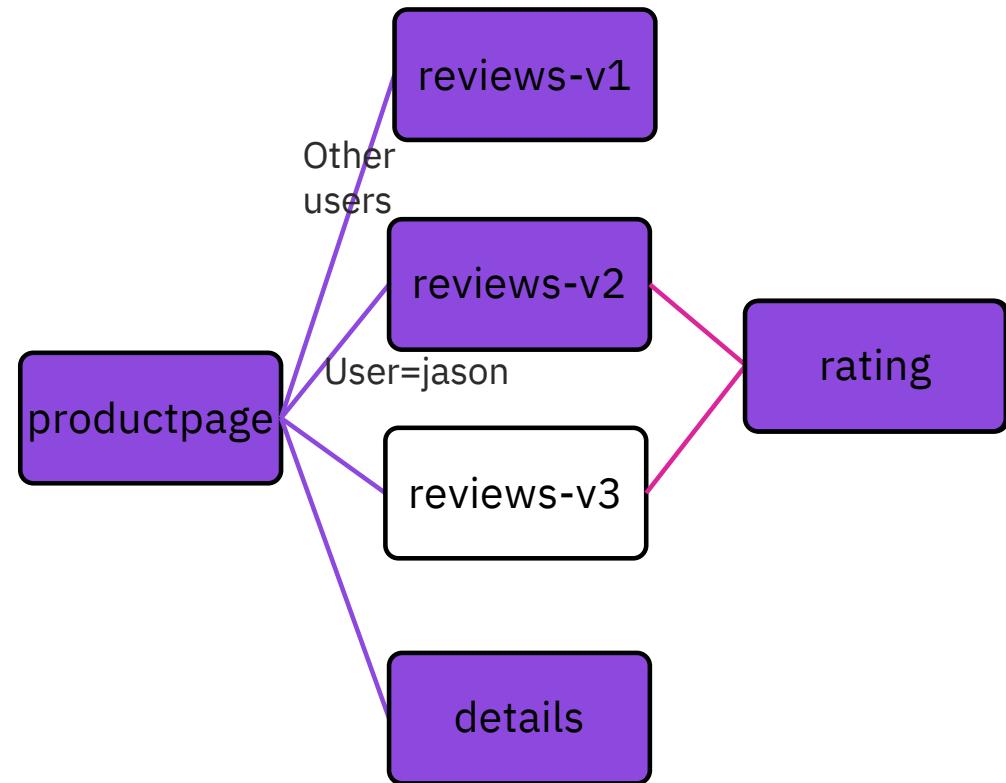


```
03route-rule-reviews-canary.yaml
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - route:
        - destination:
            host: reviews
            subset: v1
            weight: 15
        - destination:
            host: reviews
            subset: v2
            weight: 25
        - destination:
            host: reviews
            subset: v3
            weight: 60
---
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
    - name: v3
      labels:
        version: v3
```

Route based on Header

You can route to microservices based on header
(i.e. based on users, geography, browser, mobile,..
etc).

Example: Route requests based on cookie
(user=jason) to reviews-v2, else to reviews-v1.



Route based on Header

You can route to microservices based on header (i.e. based on users, geography, browser, mobile,.. etc).

Example: Route requests based on cookie (user=jason) to reviews-v2, else to reviews-v1.

```
istioctl replace -f 04route-rule-reviews-headers.yaml
```

```
04route-rule-reviews-headers.yaml
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - match:
      - headers:
          cookie:
            regex: "^(.*?;)?(user=jason)(;.*?)"$"
      route:
        - destination:
            host: reviews
            subset: v2
        - route:
            - destination:
                host: reviews
                subset: v1
```

Fault Injection

Istio enables protocol-specific fault injection into the network.

- Configure faults to be injected into requests that match specific criteria.
- Restrict the percentage of requests that should be subjected to faults.
- Two types of faults:
 - **Delay**: timing failures, mimicking increased network latency, or an overloaded upstream service.
 - **Abort**: crash failures that mimic failures in upstream services. Aborts usually manifest in the form of HTTP error codes or TCP connection failures.

Other Traffic Management Capabilities

- **HTTP Retry** describes the retry policy to use when a HTTP request fails. For example, you can define a rule that sets the maximum number of retries to 3 when calling ratings:v1 service, with a 2s timeout per retry attempt.
- **Rate Limits** allows you to deny access to services using Mixer, or specify the rate limits.
- **Circuit Breaker** allow users to set global defaults for failure recovery on a per service and/or per service version basis. It allows developers to write applications that limit the impact of failures, and latency spikes.
- **Automatic Request Timeout** allows to set timeout for services

Telemetry

Visibility

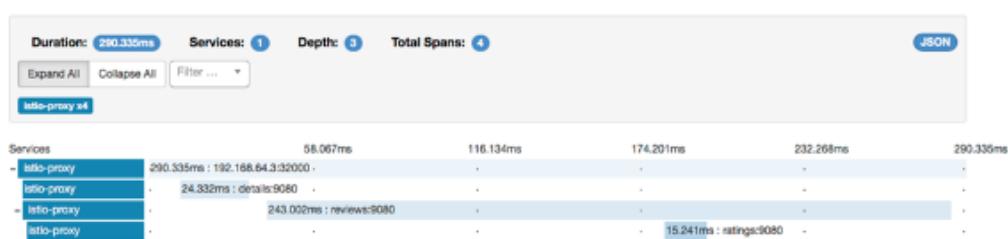
*Monitoring & tracing should not be an afterthought
in the infrastructure*

Goals

- Metrics without instrumenting apps
- Consistent metrics across fleet
- Trace flow of requests across services
- Portable across metric backend providers



Istio - Grafana dashboard w/ Prometheus backend



Istio Zipkin tracing dashboard

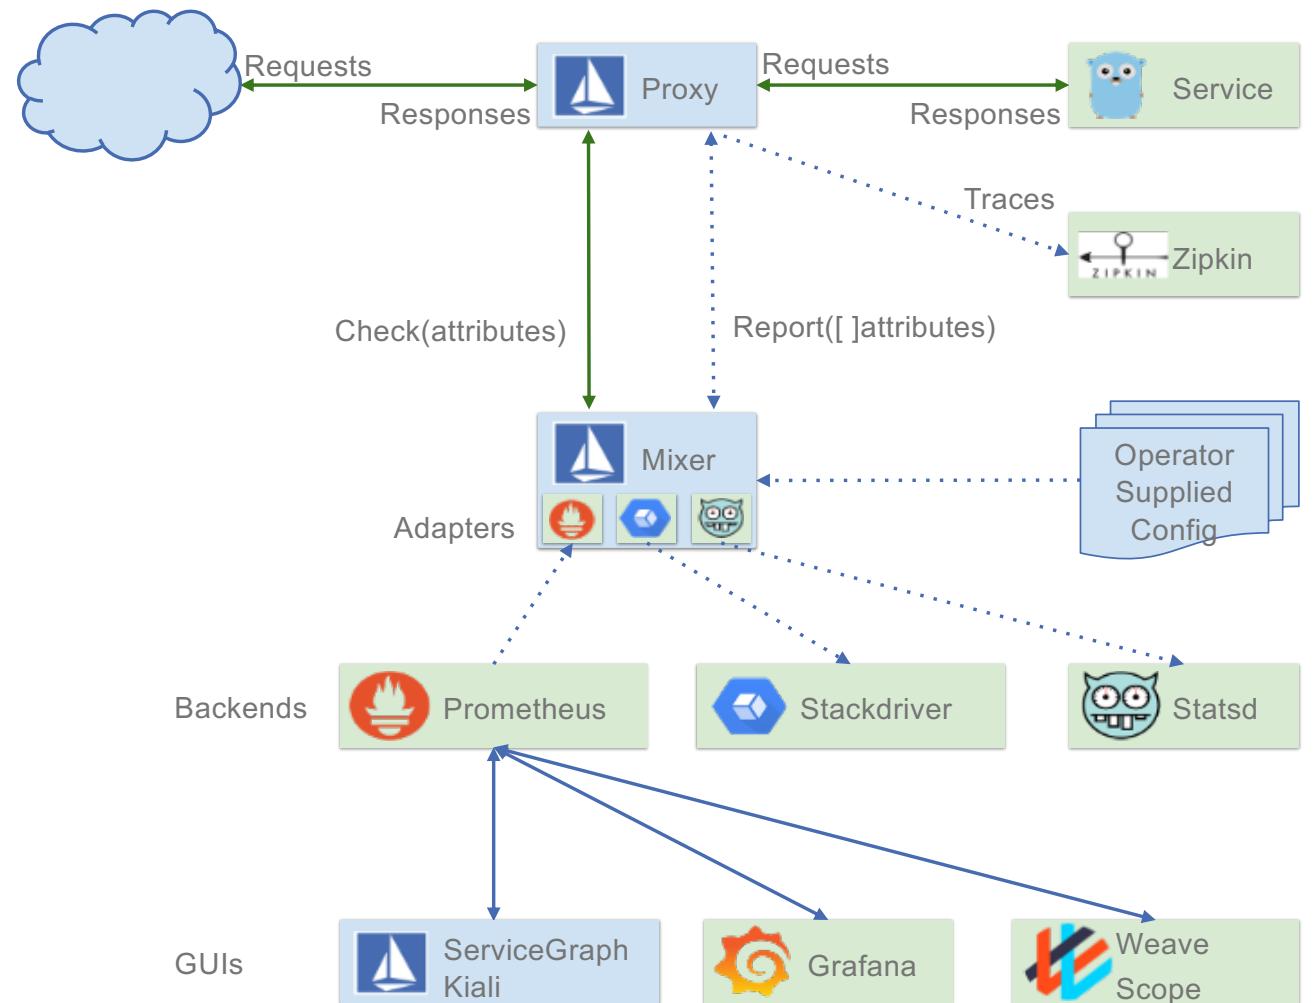
Visibility: Metrics

Mixer collects metrics emitted by Istio Proxy

Adapters in the Mixer normalize and forward to monitoring backends

Proxy send traces to Mixer

Metrics backend can be swapped at runtime



Changes needed to Legacy Applications

Modify request headers to pass span ids

- x-request-id
- x-b3-traceid
- x-b3-spanid
- x-b3-parentspanid
- x-b3-sampled
- x-b3-flags
- x-ot-span-context

Distributed Tracing: Jaeger

- Jaeger, a **CNCF** project inspired by Dapper and OpenZipkin, is a distributed tracing system released as open source by Uber Technologies.
- Jaeger is installed by default with Istio for Tracing.
- It can be used for monitoring microservices-based distributed systems:
 - Distributed context propagation
 - Distributed transaction monitoring
 - Root cause analysis
 - Service dependency analysis
 - Performance / latency optimization



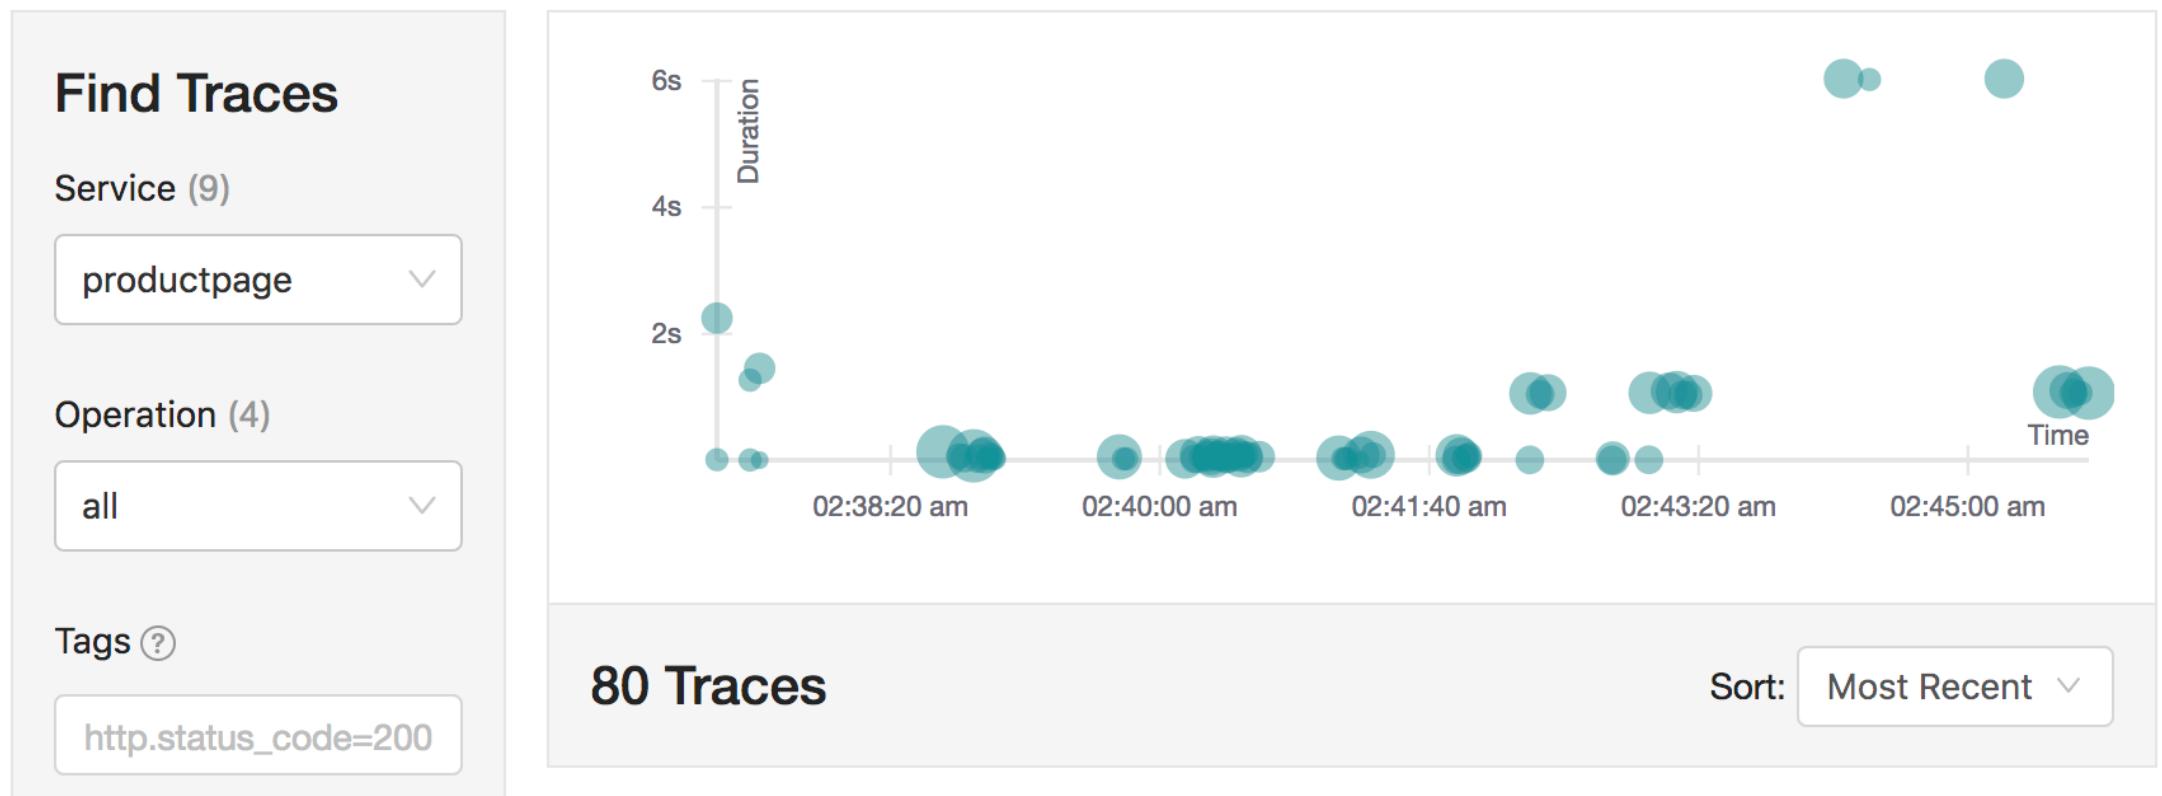
Logo by Lev Polyakov
www.polyakovproductions.com

<https://github.com/jaegertracing/jaeger>

Jaeger

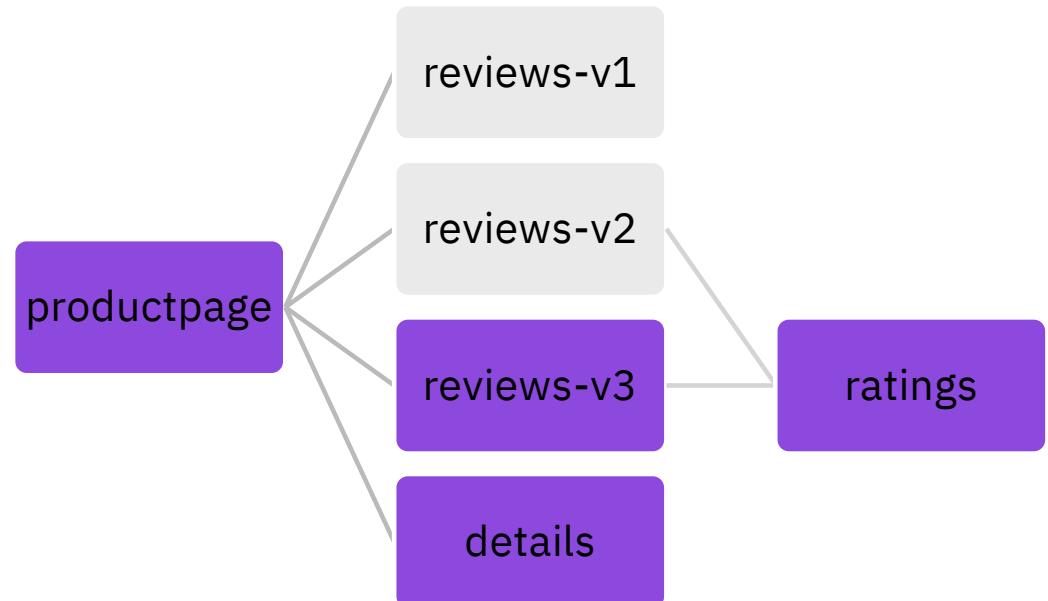
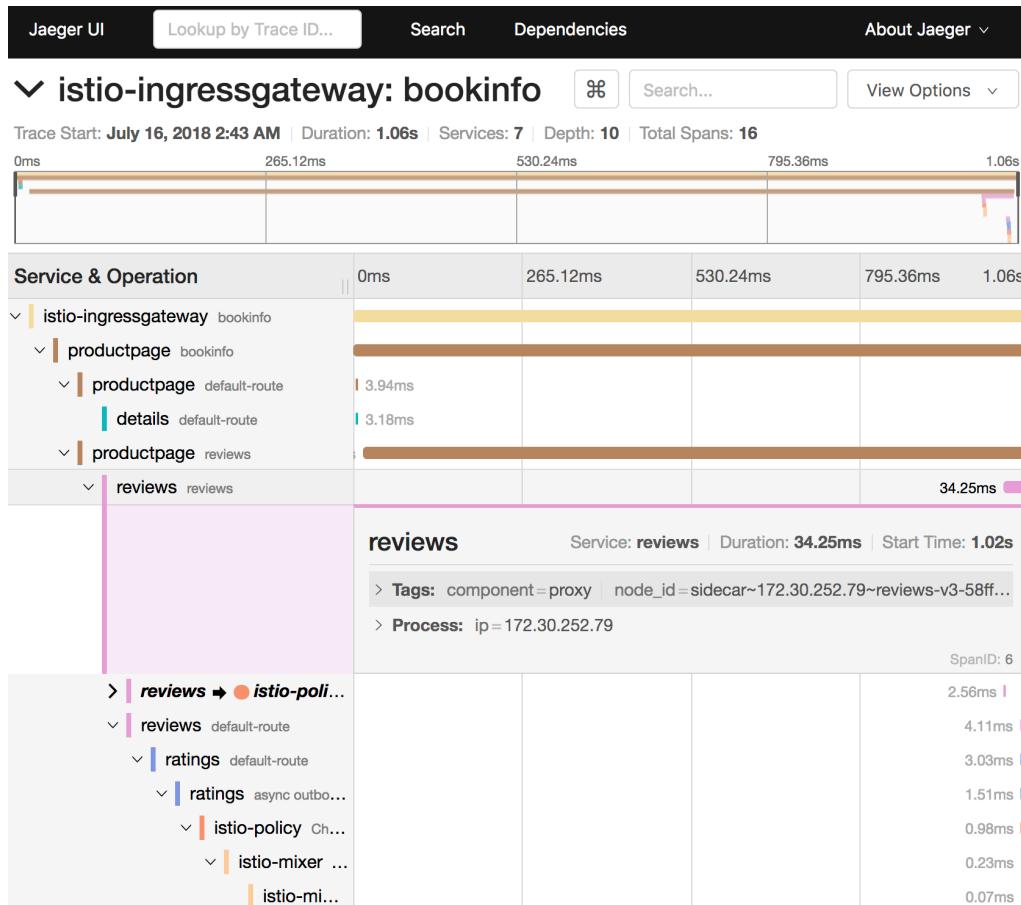
Get the tracing service details to access it.

```
kubectl get svc tracing -n istio-system
```



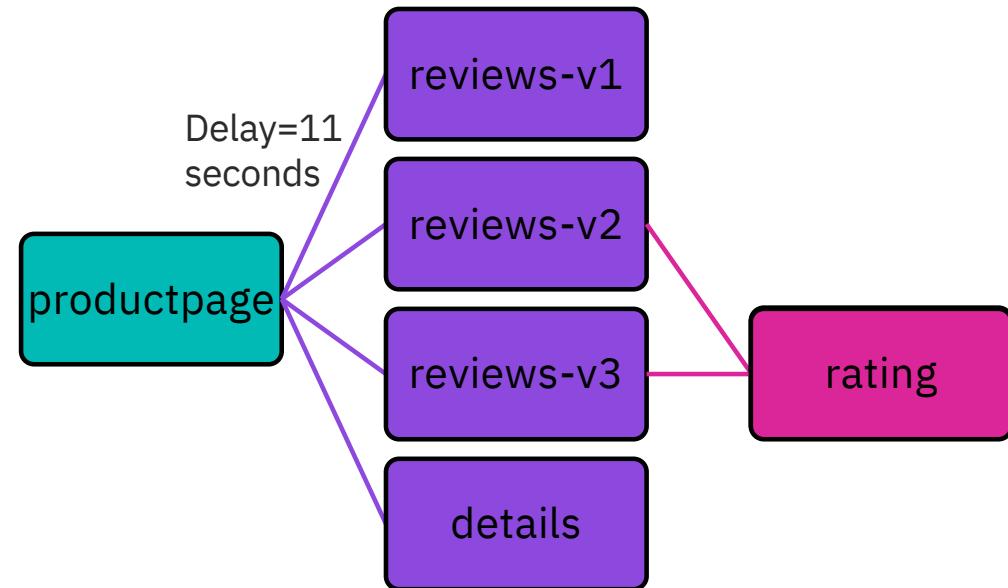
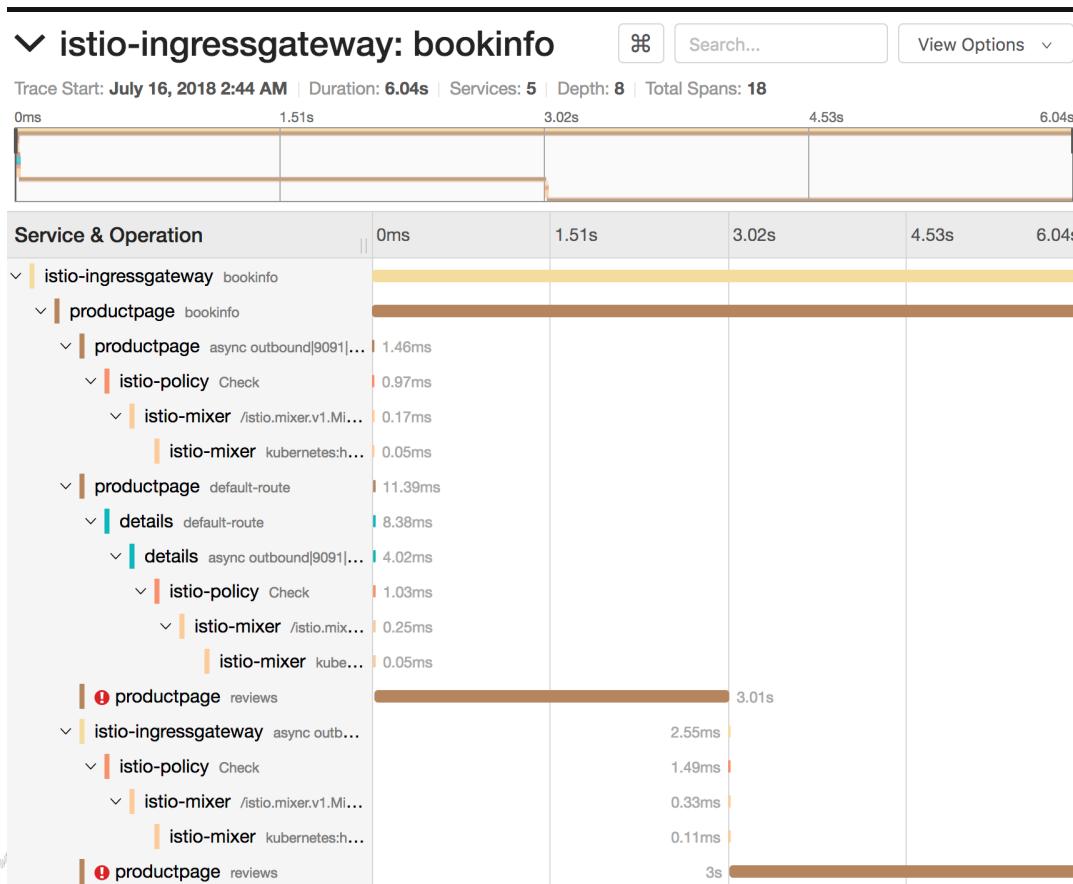
Jaeger

Inspect one of the requests



Jaeger

- Find traces for productpage where tags = error



Monitoring Visualization: Grafana

- **Grafana** is a data visualization & Monitoring with support for Graphite, InfluxDB, Prometheus, Elasticsearch and many more databases.
- It's used by default in Istio for monitoring visualization.



No matter where your data is, or what kind of database it lives in, you can bring it together with Grafana. Beautifully.

<https://grafana.com/>

Grafana

- Port forward to grafana service on Istio

```
kubectl -n istio-system port-forward $(kubectl -n istio-system get pod -l app=grafana -o jsonpath='{.items[0].metadata.name}') 3000:3000
```

- Access Grafana

<http://localhost:3000/>

- Generate some traffic

```
while sleep 3;do curl http://169.46.54.122/productpage; done
```

- The Istio Dashboard consists of three main sections:

- Global Summary View
- Mesh Summary View
- Individual Services View

Monitoring: Prometheus

- Prometheus is a CNCF project which is an open-source monitoring system with a dimensional data model, flexible query language, efficient time series database and modern alerting approach
- It's used by default in Istio for metrics monitoring.
- Istio instructs Mixer to automatically generate and report a new metric and a new log stream for all traffic within the mesh.



<https://prometheus.io/>

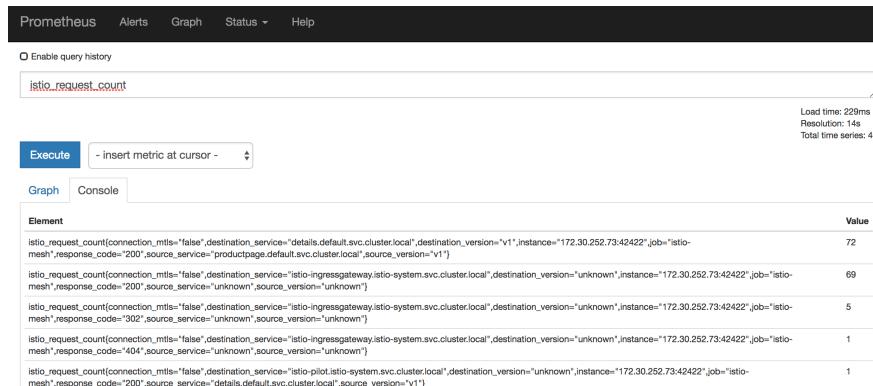
Prometheus

- Port forward to Prometheus service on Istio

```
kubectl -n istio-system port-forward $(kubectl -n istio-system get pod -l app=prometheus -o jsonpath='{.items[0].metadata.name}') 9090:9090
```

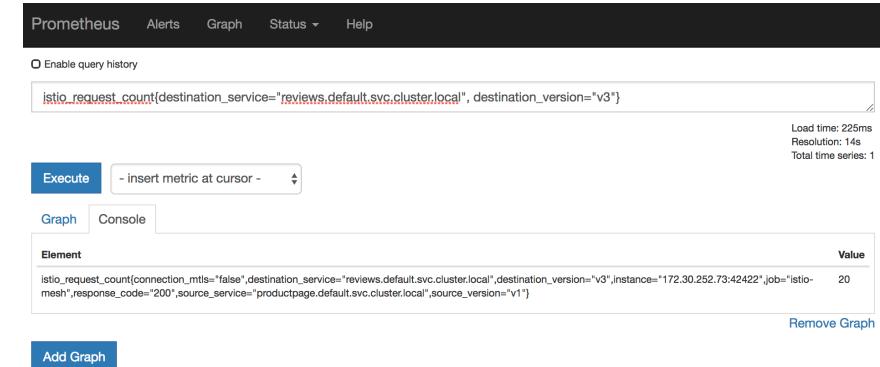
- Access Prometheus

<http://localhost:9090/graph>



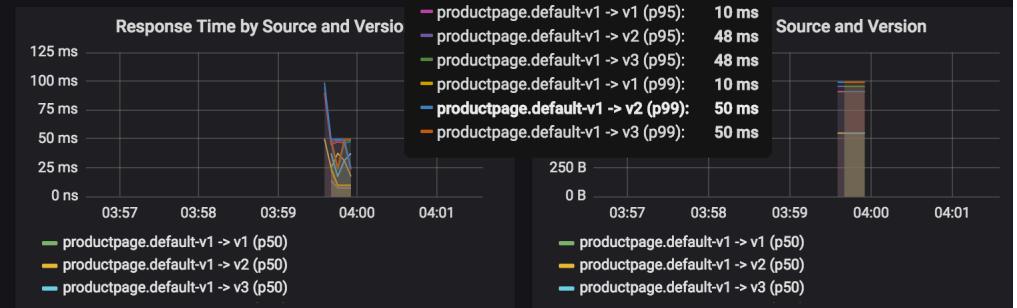
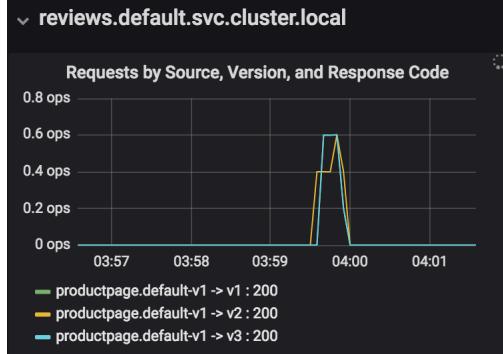
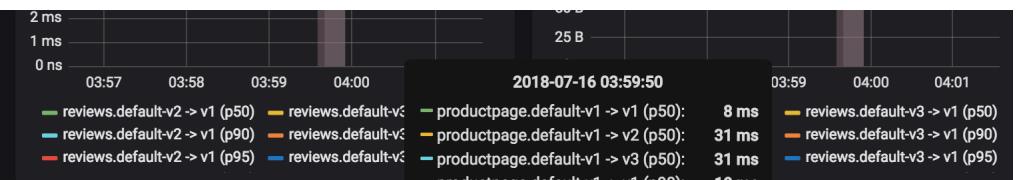
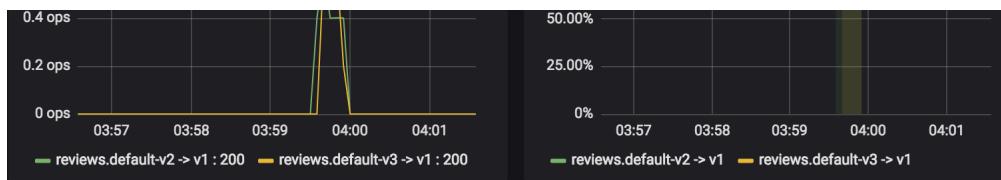
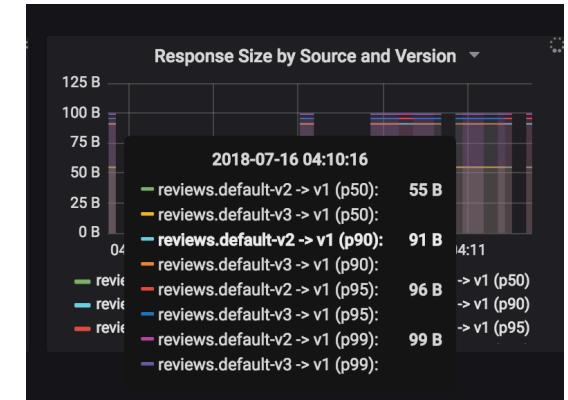
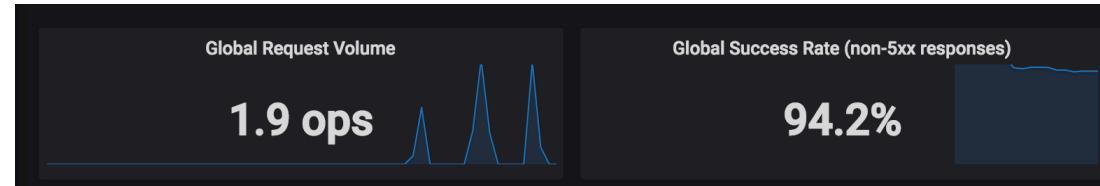
istio_request_count

IBM Developer



```
istio_request_count{  
  destination_service="reviews.default.svc.cluster.local",  
  destination_version="v3"  
}
```

Prometheus



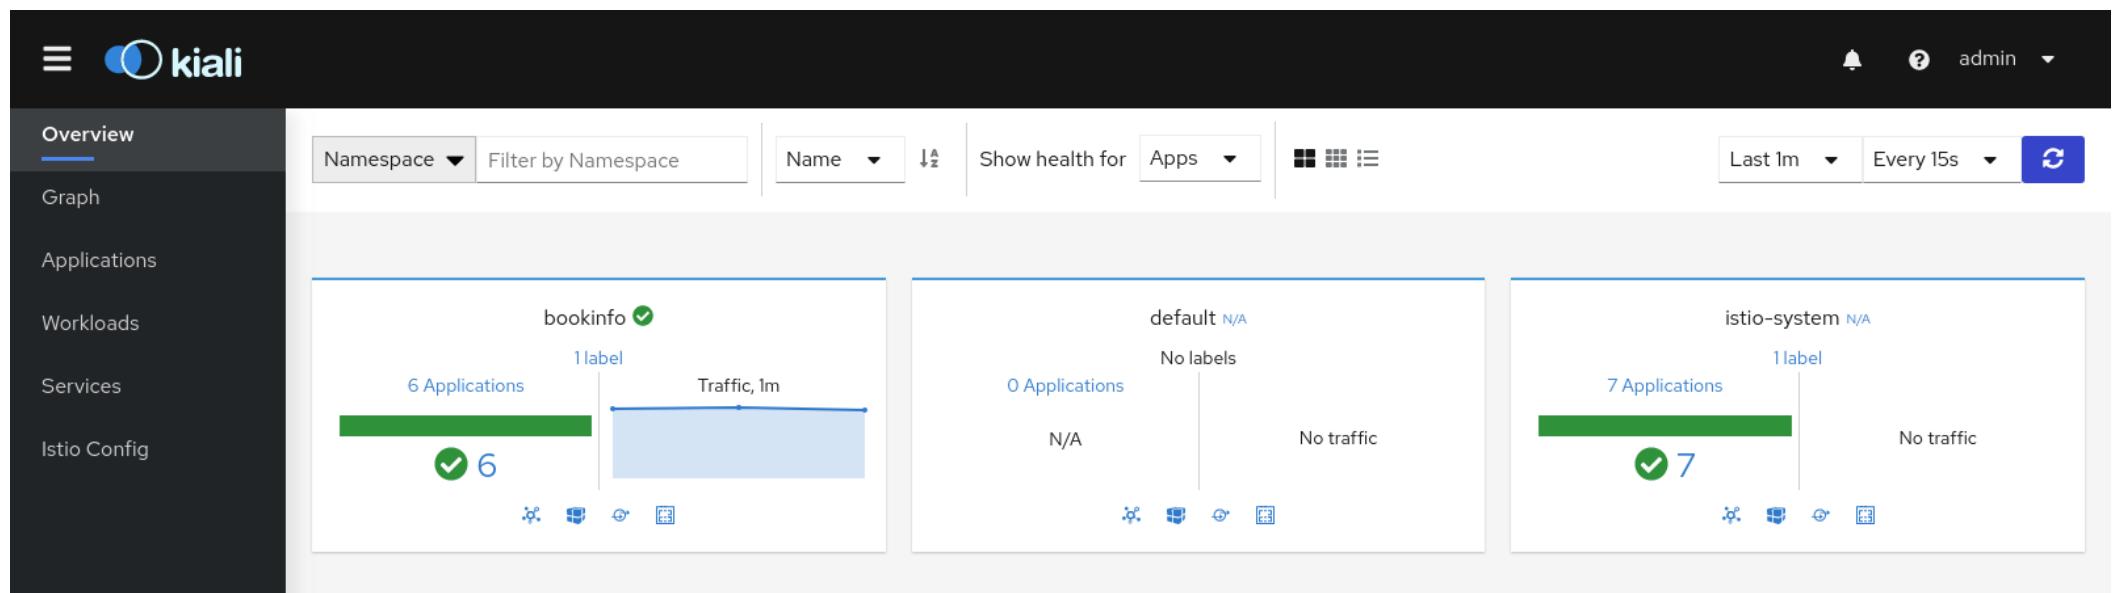
Service Graph: Kiali

- To view a graphical representation of your mesh.
- It is built on top of Prometheus queries and depends on the standard Istio metric configuration
- Port forward to service graph service on Istio

```
kubectl -n istio-system port-forward $(kubectl -n istio-system get pod -l app=kiali -o jsonpath='{.items[0].metadata.name}') 20001:20001
```

Kiali

- Access the Kiali service graph dashboard
<http://localhost:20001>



Kiali

- Generate some traffic and access application graph

Screenshot of the Kiali application interface showing the application graph and traffic metrics.

Application Graph:

- The graph displays various services and their connections. Services include `details`, `productpage`, `reviews` (v1, v2, v3), `ratings`, and `mongodb`.
- Connections are shown as directed edges between nodes. For example, `productpage` connects to `v1`, which then connects to `reviews-v1`. `reviews-v1` also connects to `v2` and `v3`.
- Annotations like `productpage` and `reviews` are placed near specific edges.

Metrics and Statistics:

- Incoming Connections:**

	Req/s
istio-ingressgateway.istio-system (unknown)	0.786441
- Outgoing Connections:**

	Req/s
details.default (v1)	0.786441
istio-pilot.istio-system (unknown)	0.000000
istio-policy.istio-system (unknown)	0.162712
istio-telemetry.istio-system (unknown)	0.681356
reviews.default (v1)	0.261017
reviews.default (v2)	0.261017
reviews.default (v3)	0.261017
- HTTP Requests per second:**

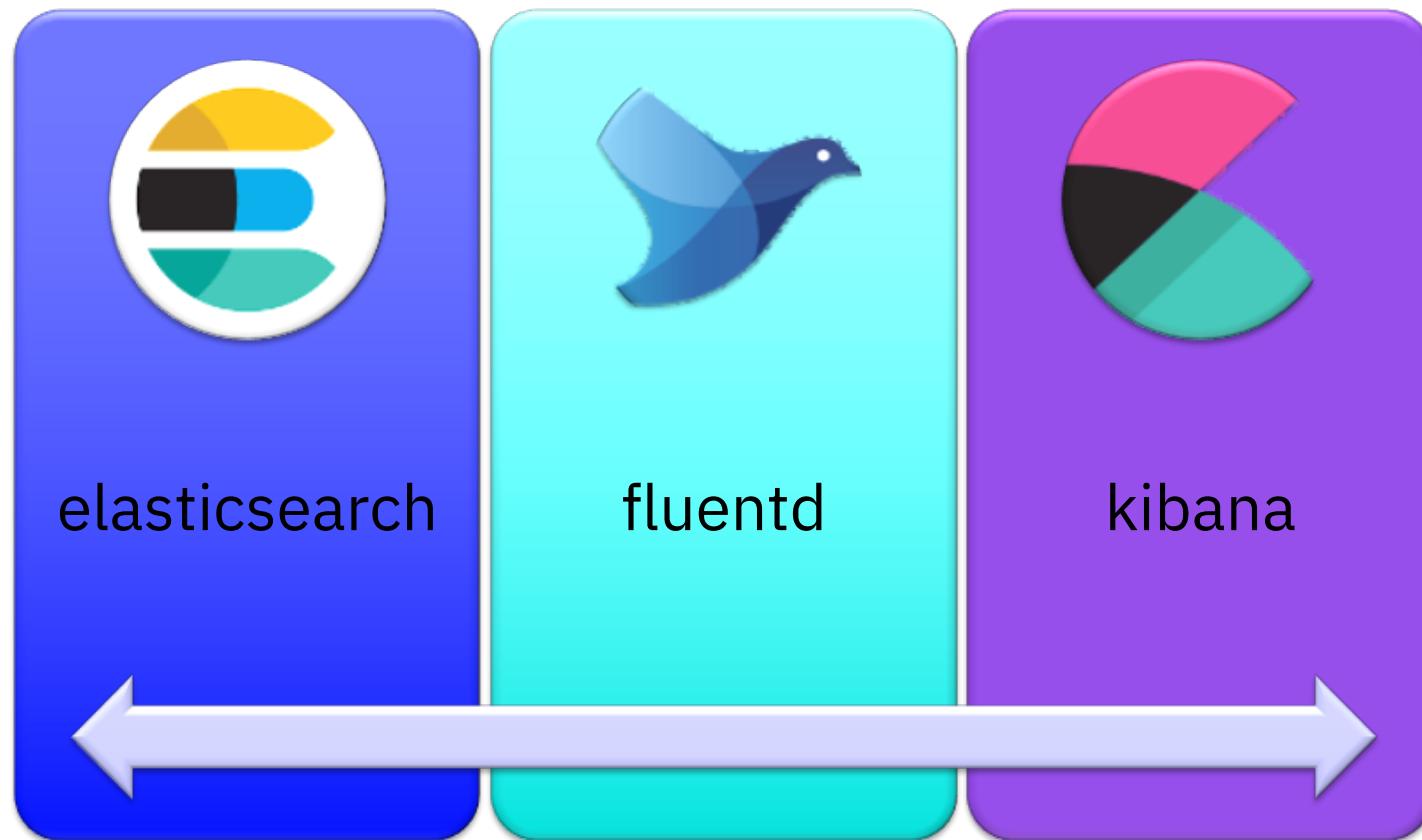
	Total	%Success	%Error
In	0.99	100.00	0.00
Out	0.64	100.00	0.00
- HTTP - Inbound Request Traffic min / max:** RPS: 0.93 / 100, %Error: 0.00 / 0.00
- HTTP - Outbound Request Traffic min / max:**

Logging

Log Aggregation

- With an increasing number of systems decoupled and scattered throughout the landscape it becomes increasingly difficult to track and trace events across all systems.
- Log aggregation solutions provides a series of benefits to distributed systems.
- The problems it tackles are:
 - Centralized, aggregated view over all log events
 - Normalization of log schema
 - Automated processing of log messages
 - Support for a great number of event sources and outputs

EFK Stack



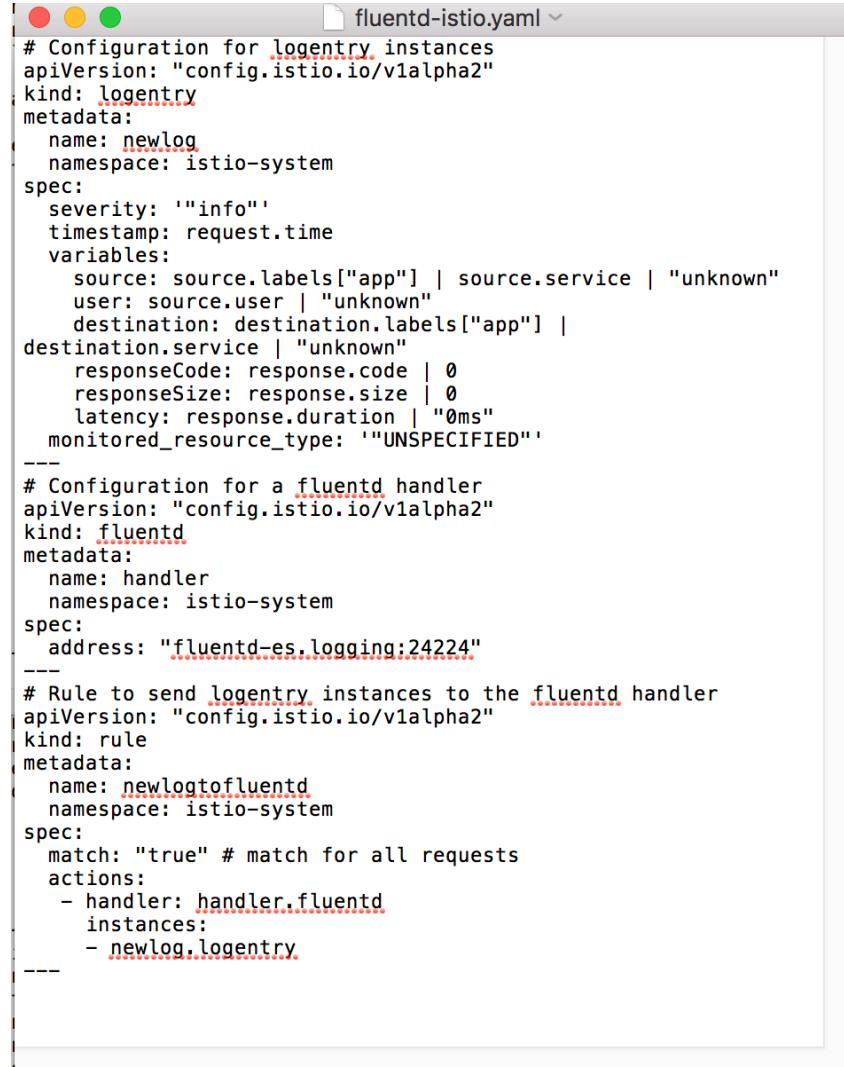
EFK Stack

- EFK stack doesn't ship by default with Istio. However, Istio allows integration with EFK stack through its proxy.
- Deploy EFK stack in your cluster.

```
kubectl apply -f logs/logging-stack.yaml
```

- Now that there is a running Fluentd daemon, add a new log entry in Istio, and send those logs to the listening daemon.
- Create a new YAML file to hold configuration for the log stream that Istio will generate and collect automatically.

```
istioctl create -f logs/fluentd-istio.yaml
```



The screenshot shows a code editor window titled "fluentd-istio.yaml". The file contains YAML configuration for Istio's logging stack. It includes three main sections: a logentry instance, a fluentd handler, and a rule. The logentry instance configures a newlog entry with specific metadata and severity levels. The fluentd handler section defines a handler named "handler" that connects to a Fluentd endpoint at port 24224. The rule section maps log entries from the logentry instance to the fluentd handler.

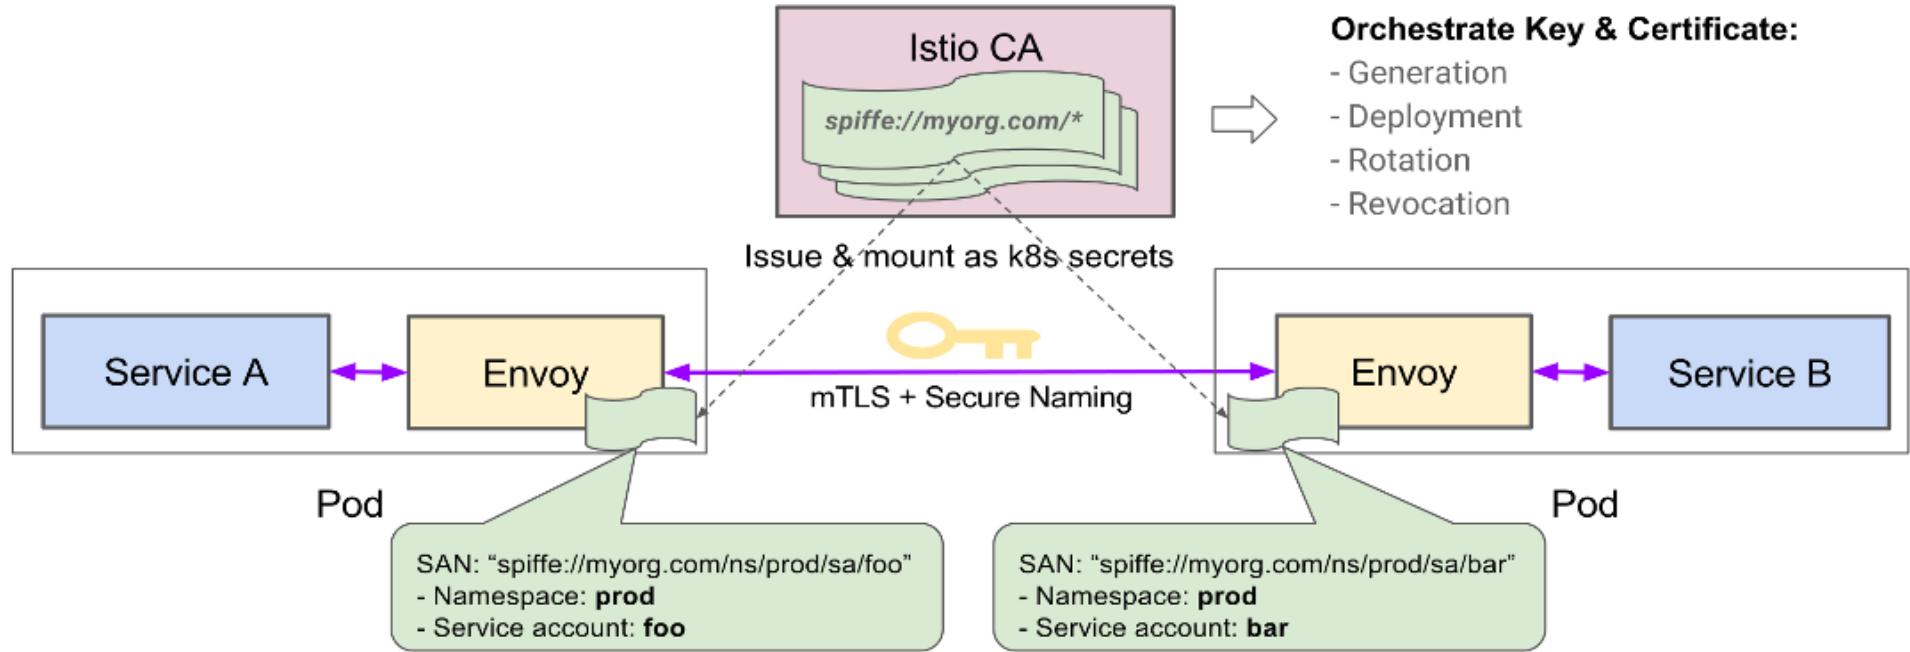
```
# Configuration for logentry instances
apiVersion: "config.istio.io/v1alpha2"
kind: logentry
metadata:
  name: newlog
  namespace: istio-system
spec:
  severity: "info"
  timestamp: request.time
variables:
  source: source.labels["app"] | source.service | "unknown"
  user: source.user | "unknown"
  destination: destination.labels["app"] | destination.service | "unknown"
  responseCode: response.code | 0
  responseSize: response.size | 0
  latency: response.duration | "0ms"
  monitored_resource_type: "'UNSPECIFIED'"
---
# Configuration for a fluentd handler
apiVersion: "config.istio.io/v1alpha2"
kind: fluentd
metadata:
  name: handler
  namespace: istio-system
spec:
  address: "fluentd-es.logging:24224"
---
# Rule to send logentry instances to the fluentd handler
apiVersion: "config.istio.io/v1alpha2"
kind: rule
metadata:
  name: newlogtofluentd
  namespace: istio-system
spec:
  match: "true" # match for all requests
  actions:
    - handler: handler.fluentd
      instances:
        - newlog.logentry
---
```

Security

Securing Microservices

- Verifiable identity
- Secure naming / addressing
- Traffic encryption
- Revocation

Istio Citadel Security at Scale



spiffe.io

Istio Citadel

- Istio can secure the communication between microservices without requiring application code changes.
- Security is provided by authenticating and encrypting communication paths within the cluster.
- Delegating communication security to Istio (as opposed to implementing TLS in each microservice), ensures that your application will be deployed with consistent and manageable security policies.
- Citadel is responsible for:
 - Providing each service with an identity representing its role.
 - Providing a common trust root to allow Envoy to validate and authenticate each other.
 - Providing a key management system

Resources

Istio documentation

<https://istio.io/docs/>

Istio in IBM products

<https://www.ibm.com/cloud/info/istio>

